

# 2025\_teaching\_pylibs

November 5, 2025

## 1 2025 Python libraries

- Presenter: Jake Krol
- Notebook contents inspired by Erik Johnson and Michael Bradshaw

### 1.1 Purpose

- Gain hands-on experience with NumPy, Pandas, Matplotlib, and other libraries

### 1.2 Methods

- Regression
- Hypothesis testing
- Gini Coefficient and Lorenz curve

### 1.3 Topics

- Climate change
- Population equilibrium
- Rice production “wealth”

### 1.4 Related work

- Lecture slides: [https://docs.google.com/presentation/d/1eq1OidZWZSsgBQHWibpLE\\_m3-NNYqYcCeVR95GhkMc/edit?usp=sharing](https://docs.google.com/presentation/d/1eq1OidZWZSsgBQHWibpLE_m3-NNYqYcCeVR95GhkMc/edit?usp=sharing)

### 1.5 Setup

- Upload `agrofood_co2_emission.csv` file

### 1.6 Scenario

You work for a data science consulting firm, and your supervisor sends you some data with minimal context. Your task is to understand what the data encodes and what insights are attainable purely from the dataset.

```
[36]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn as skl
```

```
import re
import scipy.stats as stats
```

## 1.7 Basic inspection

```
[37]: # read
# view shape
df = pd.read_csv('agrofood_co2_emission.csv')
print('# shape ', df.shape)
```

```
# shape (6965, 31)
```

```
[38]: # preview
df.head()
```

```
[38]:      Area  Year  Savanna fires  Forest fires  Crop Residues \
0  Afghanistan  1990        14.7237       0.0557    205.6077
1  Afghanistan  1991        14.7237       0.0557    209.4971
2  Afghanistan  1992        14.7237       0.0557    196.5341
3  Afghanistan  1993        14.7237       0.0557    230.8175
4  Afghanistan  1994        14.7237       0.0557    242.0494

      Rice Cultivation  Drained organic soils (CO2)  Pesticides Manufacturing \
0            686.00                  0.0          11.807483
1            678.16                  0.0          11.712073
2            686.00                  0.0          11.712073
3            686.00                  0.0          11.712073
4            705.60                  0.0          11.712073

      Food Transport  Forestland ...  Manure Management  Fires in organic soils \
0           63.1152   -2388.803 ...        319.1763             0.0
1           61.2125   -2388.803 ...        342.3079             0.0
2           53.3170   -2388.803 ...        349.1224             0.0
3           54.3617   -2388.803 ...        352.2947             0.0
4           53.9874   -2388.803 ...        367.6784             0.0

      Fires in humid tropical forests  On-farm energy use  Rural population \
0                      0.0                 NaN     9655167.0
1                      0.0                 NaN    10230490.0
2                      0.0                 NaN    10995568.0
3                      0.0                 NaN    11858090.0
4                      0.0                 NaN    12690115.0

      Urban population  Total Population - Male  Total Population - Female \
0           2593947.0            5348387.0            5346409.0
1           2763167.0            5372959.0            5372208.0
2           2985663.0            6028494.0            6028939.0
```

```
3      3237009.0      7003641.0      7000119.0
4      3482604.0      7733458.0      7722096.0
```

```
total_emission  Average Temperature °C
0      2198.963539      0.536167
1      2323.876629      0.020667
2      2356.304229     -0.259583
3      2368.470529      0.101917
4      2500.768729      0.372250
```

[5 rows x 31 columns]

```
[39]: # count column data types
df.dtypes.value_counts()
```

```
[39]: float64    29
object       1
int64       1
Name: count, dtype: int64
```

checkpoint - Agricultural data - For various (area,year) combinations, there exists numeric features - e.g., Afghan 1990 rice cultivation questions - How many distinct countries exist in df? - What is the year span? - What are the units of measurement for these columns? - What are mean/expected values in each numeric column?

```
[40]: # distinct areas
print(df['Area'].nunique())
print(df['Area'].unique()[:10])
```

236  
['Afghanistan' 'Albania' 'Algeria' 'American Samoa' 'Andorra' 'Angola'  
'Anguilla' 'Antigua and Barbuda' 'Argentina' 'Armenia']

```
[41]: # year span
print(
    df['Year'].min(),
    df['Year'].max()
)
```

1990 2020

```
[42]: df[df.columns.drop(['Area', 'Year'])].describe().T['mean'].round(2)
```

```
[42]: Savanna fires           1188.39
Forest fires            919.30
Crop Residues           998.71
Rice Cultivation        4259.67
Drained organic soils (CO2) 3503.23
Pesticides Manufacturing   333.42
```

Food Transport	1939.58
Forestland	-17828.29
Net Forest conversion	17605.64
Food Household Consumption	4847.58
Food Retail	2043.21
On-farm Electricity Use	1626.68
Food Packaging	1658.63
Agrifood Systems Waste Disposal	6018.44
Food Processing	3872.72
Fertilizers Manufacturing	3035.72
IPPU	19991.50
Manure applied to Soils	923.23
Manure left on Pasture	3518.03
Manure Management	2263.34
Fires in organic soils	1210.32
Fires in humid tropical forests	668.45
On-farm energy use	3008.98
Rural population	17857735.39
Urban population	16932296.97
Total Population - Male	17619629.63
Total Population - Female	17324469.29
total_emission	64091.24
Average Temperature °C	0.87

Name: mean, dtype: float64

questions - why is forestland negative? is change in forestland measured? - given this data, what hypotheses could we form? - let's get a frame of reference by inspecting one area/country

```
[43]: df_usa=df[df['Area'] == 'United States of America']
df_usa.head()
```

	Area	Year	Savanna fires	Forest fires	\
6591	United States of America	1990	1391.1481	1999.5617	
6592	United States of America	1991	1391.1481	1999.5617	
6593	United States of America	1992	1391.1481	1999.5617	
6594	United States of America	1993	1391.1481	1999.5617	
6595	United States of America	1994	1391.1481	1999.5617	
	Crop Residues	Rice Cultivation	Drained organic soils (CO2)	\	
6591	19193.1314	11195.520		50713.3546	
6592	17574.1126	11029.312		50713.3546	
6593	20964.8025	12421.402		50713.3546	
6594	16761.5327	11235.602		50566.2744	
6595	21370.6468	13151.110		50418.1314	
	Pesticides Manufacturing	Food Transport	Forestland	...	\
6591		13772.0	45410.2673	-520573.1593	...
6592		13139.0	44566.9408	-520573.1593	...

```

6593          14109.0      46198.5933 -520573.1593 ...
6594          12948.0      45175.6740 -520573.1593 ...
6595          14401.0      46739.0889 -520573.1593 ...

    Manure Management  Fires in organic soils \
6591          47024.7024      0.0
6592          48206.0451      0.0
6593          48416.2439      0.0
6594          48426.0899      0.0
6595          49165.7227      0.0

    Fires in humid tropical forests  On-farm energy use  Rural population \
6591          9.8513        54454.7092       62373717.0
6592          9.8513        54565.6091       61957131.0
6593          9.8513        57469.9273       61539241.0
6594          9.8513        54689.3180       61136396.0
6595          9.8513        55539.8937       60759735.0

    Urban population  Total Population - Male  Total Population - Female \
6591          190156233.0      121451448.0       126632284.0
6592          193017688.0      123229931.0       128330258.0
6593          195915032.0      125081499.0       130093840.0
6594          198883790.0      126914111.0       131865642.0
6595          201981831.0      128685441.0       133588147.0

    total_emission  Average Temperature °C
6591          463050.9394      0.733583
6592          473285.7816      0.706333
6593          486026.3425      0.253000
6594          484238.6197      0.153500
6595          509412.4984      0.470250

```

[5 rows x 31 columns]

questions - we have savannas in the USA? - does avg. temp support climate warming hypotheses?  
- are urban/rural populations shrinking/growing?

```
[44]: # graph avg temp over time in the US
# plt.plot(df_usa['Year'], df_usa['AvgTemp'])
# plt.show()
```

- oh, that's very annoying. the temperature column uses the `°` character ...
- what's a simple way to resolve this?

```
[45]: # let's clean up column names
def cln(x):
    x = x.replace('°', '')
    # whitespace
```

```

x = re.sub(r"\s+", " ", x)
# hyphens
x = x.replace('-', '_')
# parentheses
x = x.replace('(', '')
x = x.replace(')', '')
# redundant '_' removal
x = re.sub(r"_+", " ", x)
# lower
return x.lower()
df.columns = [cln(x) for x in df.columns]
df.head()

```

[45]:

	area	year	savanna_fires	forest_fires	crop_residues	\
0	Afghanistan	1990	14.7237	0.0557	205.6077	
1	Afghanistan	1991	14.7237	0.0557	209.4971	
2	Afghanistan	1992	14.7237	0.0557	196.5341	
3	Afghanistan	1993	14.7237	0.0557	230.8175	
4	Afghanistan	1994	14.7237	0.0557	242.0494	
	rice_cultivation	drained_organic_soils_co2	pesticides_manufacturing	\		
0	686.00		0.0		11.807483	
1	678.16		0.0		11.712073	
2	686.00		0.0		11.712073	
3	686.00		0.0		11.712073	
4	705.60		0.0		11.712073	
	food_transport	forestland	...	manure_management	fires_in_organic_soils	\
0	63.1152	-2388.803	...	319.1763		0.0
1	61.2125	-2388.803	...	342.3079		0.0
2	53.3170	-2388.803	...	349.1224		0.0
3	54.3617	-2388.803	...	352.2947		0.0
4	53.9874	-2388.803	...	367.6784		0.0
	fires_in_humid_tropical_forests	on_farm_energy_use	rural_population	\		
0		0.0	NaN		9655167.0	
1		0.0	NaN		10230490.0	
2		0.0	NaN		10995568.0	
3		0.0	NaN		11858090.0	
4		0.0	NaN		12690115.0	
	urban_population	total_population_male	total_population_female	\		
0	2593947.0	5348387.0		5346409.0		
1	2763167.0	5372959.0		5372208.0		
2	2985663.0	6028494.0		6028939.0		
3	3237009.0	7003641.0		7000119.0		
4	3482604.0	7733458.0		7722096.0		

```

total_emission  average_temperature_c
0      2198.963539          0.536167
1      2323.876629          0.020667
2      2356.304229         -0.259583
3      2368.470529          0.101917
4      2500.768729          0.372250

```

[5 rows x 31 columns]

```
[46]: # clean usa subset
df_usa.columns = [cln(x) for x in df_usa.columns]
df_usa.head()
```

```

[46]:           area  year  savanna_fires  forest_fires \
6591  United States of America  1990      1391.1481    1999.5617
6592  United States of America  1991      1391.1481    1999.5617
6593  United States of America  1992      1391.1481    1999.5617
6594  United States of America  1993      1391.1481    1999.5617
6595  United States of America  1994      1391.1481    1999.5617

           crop_residues  rice_cultivation  drained_organic_soils_co2 \
6591      19193.1314          11195.520          50713.3546
6592      17574.1126          11029.312          50713.3546
6593      20964.8025          12421.402          50713.3546
6594      16761.5327          11235.602          50566.2744
6595      21370.6468          13151.110          50418.1314

           pesticides_manufacturing  food_transport  forestland ... \
6591                  13772.0        45410.2673 -520573.1593 ...
6592                  13139.0        44566.9408 -520573.1593 ...
6593                  14109.0        46198.5933 -520573.1593 ...
6594                  12948.0        45175.6740 -520573.1593 ...
6595                  14401.0        46739.0889 -520573.1593 ...

           manure_management  fires_in_organic_soils \
6591                  47024.7024          0.0
6592                  48206.0451          0.0
6593                  48416.2439          0.0
6594                  48426.0899          0.0
6595                  49165.7227          0.0

           fires_in_humid_tropical_forests  on_farm_energy_use  rural_population \
6591                      9.8513        54454.7092       62373717.0
6592                      9.8513        54565.6091       61957131.0
6593                      9.8513        57469.9273       61539241.0
6594                      9.8513        54689.3180       61136396.0

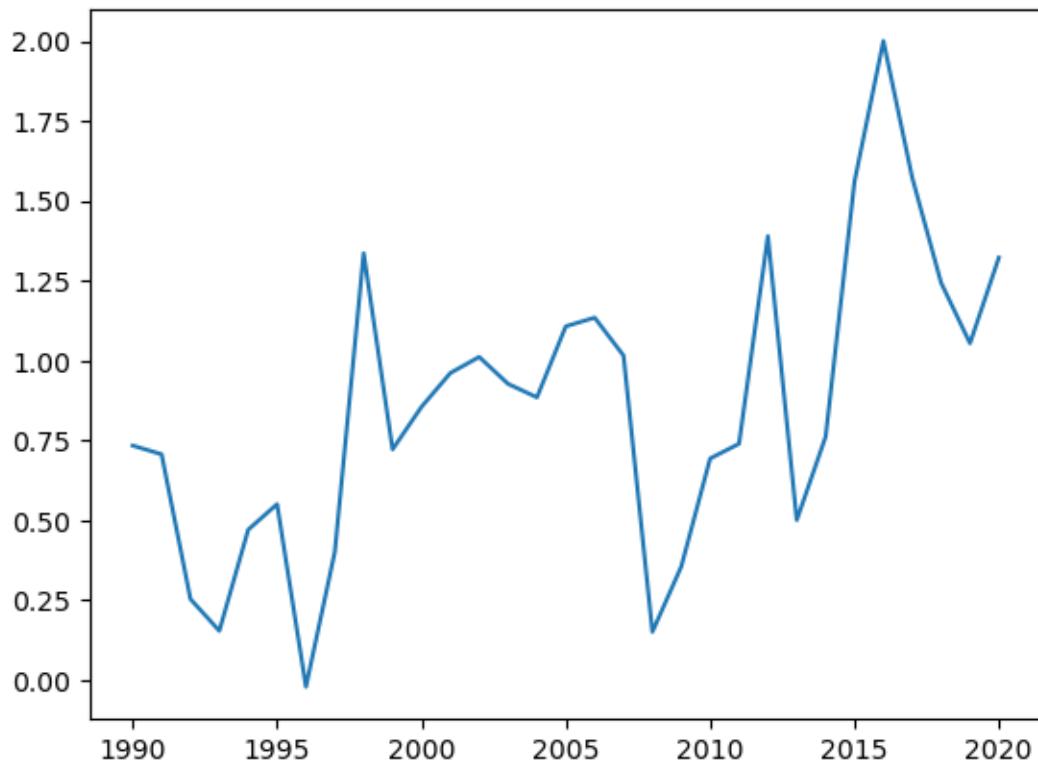
```

6595		9.8513	55539.8937	60759735.0
	urban_population	total_population_male	total_population_female	\
6591	190156233.0	121451448.0	126632284.0	
6592	193017688.0	123229931.0	128330258.0	
6593	195915032.0	125081499.0	130093840.0	
6594	198883790.0	126914111.0	131865642.0	
6595	201981831.0	128685441.0	133588147.0	
	total_emission	average_temperature_c		
6591	463050.9394	0.733583		
6592	473285.7816	0.706333		
6593	486026.3425	0.253000		
6594	484238.6197	0.153500		
6595	509412.4984	0.470250		

[5 rows x 31 columns]

```
[47]: # is climate warming hypothesis supported by USA data?
x = df_usa['year']
y = df_usa['average_temperature_c']
plt.plot(x,y)
```

[47]: [`<matplotlib.lines.Line2D at 0x7f8171643d10>`]



interesting, there's an upward trend.

Q: how could we simply, numerically quantify this trend?

Click to reveal spoiler

linear regression

## 1.8 Climate change regression

```
[48]: # regression with OLS objective
# using scipy
slope, intercept, r_value, p_value, std_err = stats.linregress(
    x,
    y,
    alternative='greater'
)
print(
    'slope: ', slope,
    '\nintercept: ', intercept,
    '\nr_value: ', r_value,
    '\np_value: ', p_value,
    '\nstd_err: ', std_err
)
```

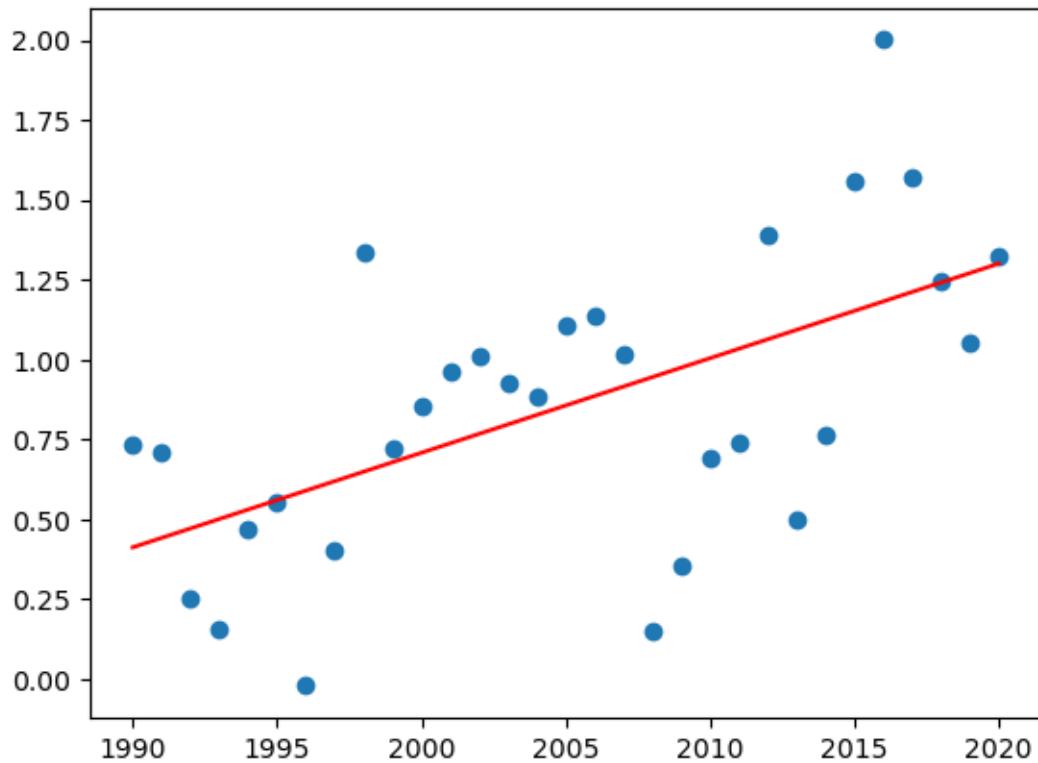
```
slope: 0.02962479838709677
intercept: -58.5416320564516
r_value: 0.5784210865537276
p_value: 0.0003266897078220169
std_err: 0.007758239185112271
```

checkpoint: regression analysis - how can we visualize our regression line - why did we get a p-value? how do we interpret it in this context? (hint: what did we talk about in monte carlo hypothesis testing?)

Click to reveal spoiler

- slope and intercept provide the linear equation; therefore, we can plug in arbitrary inputs in the domain  $x \in R$ . we're interested

```
[49]: # plot fit regression line and data
plt.plot(x, y, 'o')
plt.plot(x, intercept + slope*x, 'r')
plt.show()
```



Yes, from the data, we can **reject the null hypothesis** ( $p = .00032$ ) that average temperature is non-increasing over time ... in the USA.

What about other countries? How would you design and implement this experiment on a larger scale?

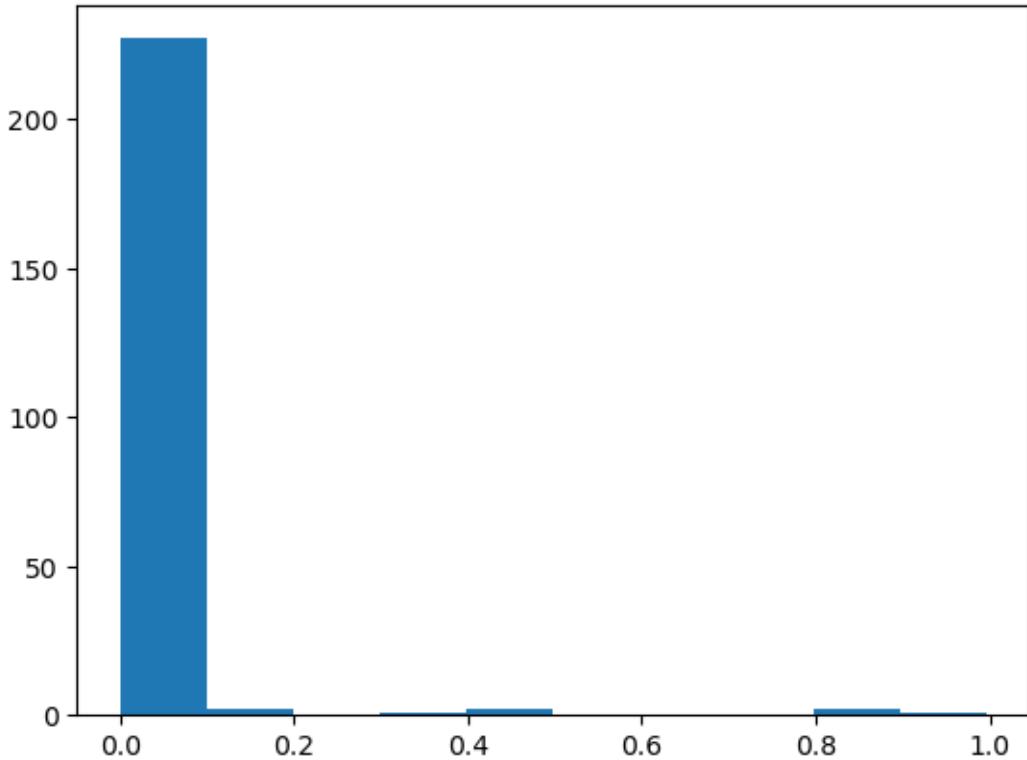
```
[50]: def gwarm_exp(years, avg_temps):
    slope, _, _, p_value, _ = stats.linregress(
        years,
        avg_temps,
        alternative='greater'
    )
    return p_value, slope
# apply this experiment to all areas
p_values = []
slopes = []
for group, df_group in df.groupby('area'):
    p_value, slope = gwarm_exp(df_group['year'],
                                df_group['average_temperature_c'])
    p_values.append(p_value)
    slopes.append(slope)
# plot distribution of p-values
plt.hist(p_values)
```

```

/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10729: RuntimeWarning: invalid value
encountered in scalar divide
    slope = ssxym / ssxm
/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10743: RuntimeWarning: invalid value
encountered in sqrt
    t = r * np.sqrt(df / ((1.0 - r + TINY)*(1.0 + r + TINY)))
/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10749: RuntimeWarning: invalid value
encountered in scalar divide
    slope_stderr = np.sqrt((1 - r**2) * ssym / ssxm / df)

[50]: (array([227.,  2.,  0.,  1.,  2.,  0.,  0.,  0.,  2.,  1.]),
       array([0.          , 0.09962106, 0.19924212, 0.29886318, 0.39848424,
              0.49810531, 0.59772637, 0.69734743, 0.79696849, 0.89658955,
              0.99621061]),
       <BarContainer object of 10 artists>

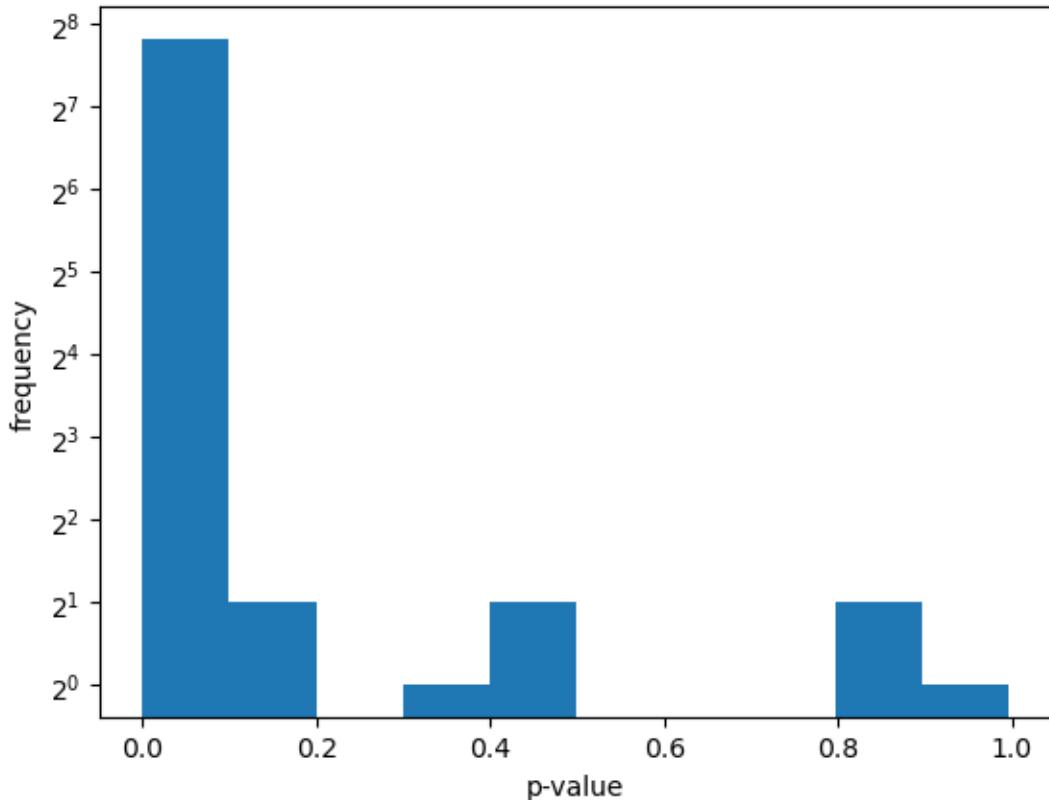
```



most p-values are close to zero. But,  $\exists$  some  $p \gg 0$ .

How can we make this plot better?

```
[51]: plt.hist(p_values)
plt.yscale("log", base = 2)
plt.xlabel("p-value")
plt.ylabel("frequency")
plt.show()
```



now, high p-values can easily be counted (cases where climate warming hyp was not supported).

i counted 8.

let's be more precise

```
[52]: # what fraction of areas support climate warming hypothesis
# try diff alpha and see how fraction changes
alpha = 0.05
n = len(p_values)
num = [p < alpha for p in p_values].count(True)
print(num/n)
```

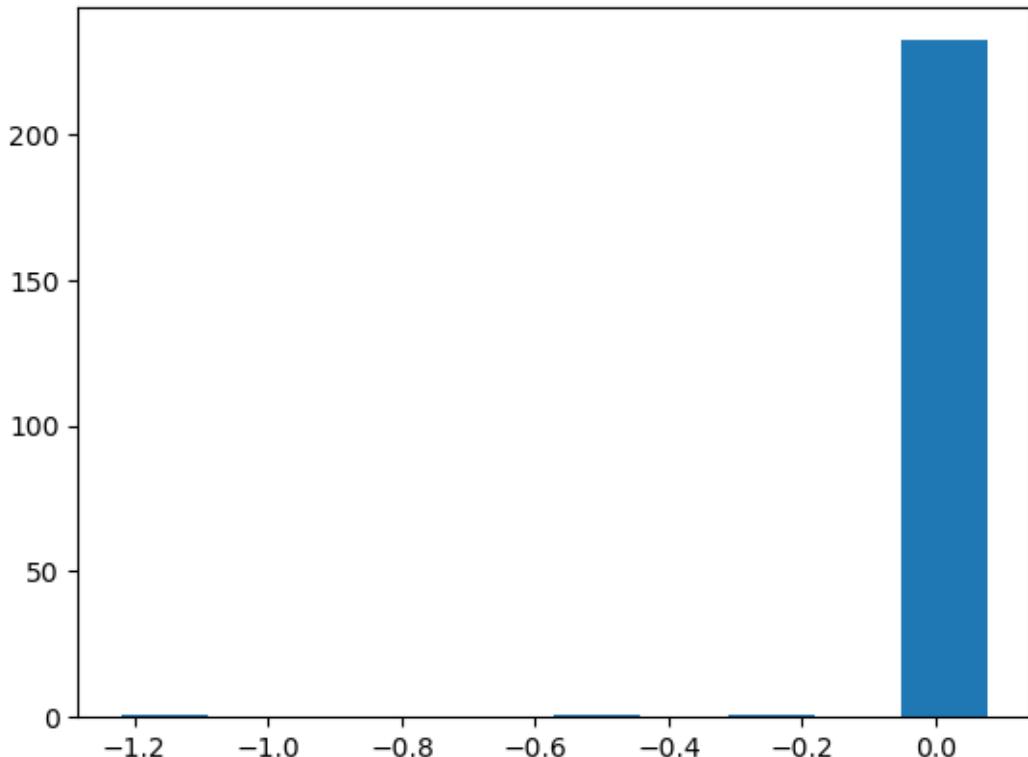
0.940677966101695

conclusion: 94% of areas in our dataset support a climate warming hypothesis

we showed that temperature is rising, but to what degree and in which places?

```
[53]: plt.hist(slopes)
```

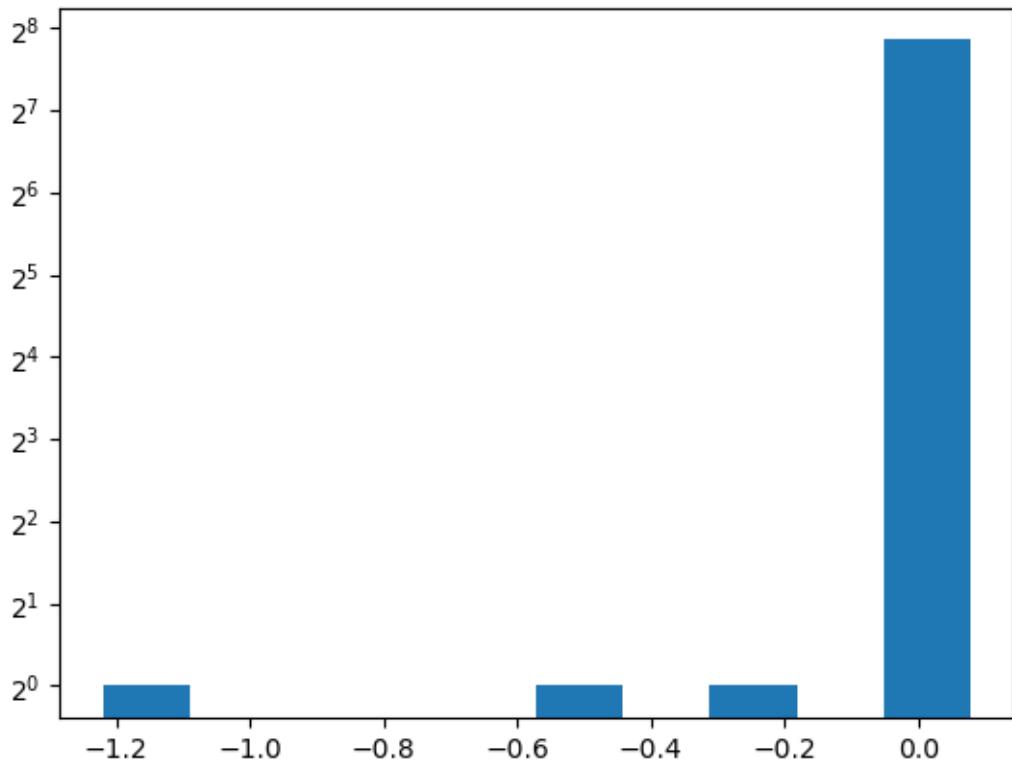
```
[53]: (array([ 1.,  0.,  0.,  0.,  1.,  0.,  1.,  0., 232.]),
array([-1.21991667, -1.09021346, -0.96051025, -0.83080704, -0.70110383,
-0.57140062, -0.44169741, -0.3119942 , -0.18229099, -0.05258778,
0.07711544]),  
<BarContainer object of 10 artists>)
```



clearly many slopes are close to zero, ... however, the tail of this distribution is interesting.

some places have a strong decrease in temp over the years

```
[54]: # again we scale y-axis for better view of dist. tail
plt.hist(slopes)
plt.yscale("log", base = 2)
plt.show()
```



two paths 1. i want to show evidence against climate warming 2. i want to show evidence supporting climate warming

```
[55]: # path 1
# goal: find outlier areas where avg temp is decreasing sharply
p_values_cherry = []
slopes_cherry = []
groups_cherry = []
for group, df_group in df.groupby('area'):
    p_value, slope = gwarm_exp(df_group['year'], df_group['average_temperature_c'])
    if slope < -0.1:
        p_values_cherry.append(p_value)
        slopes_cherry.append(slope)
        groups_cherry.append(group)
groups_cherry
```

```
/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10729: RuntimeWarning: invalid value
encountered in scalar divide
    slope = ssxym / ssxm
/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10743: RuntimeWarning: invalid value
```

```

encountered in sqrt
t = r * np.sqrt(df / ((1.0 - r + TINY)*(1.0 + r + TINY)))
/data/jake/miniconda3/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:10749: RuntimeWarning: invalid value
encountered in scalar divide
slope_stderr = np.sqrt((1 - r**2) * ssym / ssxm / df)

```

[55]: ['Ethiopia PDR', 'USSR', 'Yugoslav SFR']

[56]: # path 1 cont.

```

df_cherry = df[df['area'].isin(groups_cherry)]
df_cherry.head()

```

	area	year	savanna_fires	forest_fires	crop_residues	\
2152	Ethiopia PDR	1990	3851.3866	1509.6419	349.6276	
2153	Ethiopia PDR	1991	3851.3866	1509.6419	286.8652	
2154	Ethiopia PDR	1992	3851.3866	1509.6419	270.0245	
6684	USSR	1990	8405.2264	7262.4148	14854.7660	
6685	USSR	1991	8405.2264	7262.4148	11900.4386	
	rice_cultivation	drained_organic_soils_co2	pesticides_manufacturing	\		
2152	3391.507802		5513.0102		12.850182	
2153	3309.287814		5513.0102		12.098396	
2154	3586.696796		5513.0102		11.606165	
6684	4813.760000		131838.2352		1169.000000	
6685	4699.296000		131838.2352		1301.000000	
	food_transport	forestland	...	manure_management	\	
2152	100.5070	-107.2750	...	1543.7229		
2153	98.8585	-107.2750	...	1554.3704		
2154	33.8336	-107.2750	...	1594.3863		
6684	32821.6716	-605722.9991	...	60407.5004		
6685	32500.3786	-605722.9991	...	59296.5706		
	fires_in_organic_soils	fires_in_humid_tropical_forests	\			
2152	0.0		416.3368			
2153	0.0		416.3368			
2154	0.0		416.3368			
6684	0.0		0.0000			
6685	0.0		0.0000			
	on_farm_energy_use	rural_population	urban_population	\		
2152	80.8401	44542203.0	6657624.0			
2153	82.1289	45925516.0	7022864.0			
2154	41.6932	47359180.0	7407170.0			
6684	248879.1769	99158922.0	188910774.0			
6685	246785.4967	99725999.0	189602194.0			

```

      total_population_male  total_population_female  total_emission \
2152          24890395.0            25137639.0    65914.320714
2153          25886400.0            26090250.0    61115.468490
2154          26890139.0            27042568.0    61654.073430
6684          136777703.0           153126995.0   524473.945541
6685          137641632.0           153853429.0   520480.833615

      average_temperature_c
2152          0.625000
2153          0.342667
2154          0.048444
6684          1.158250
6685          0.609000

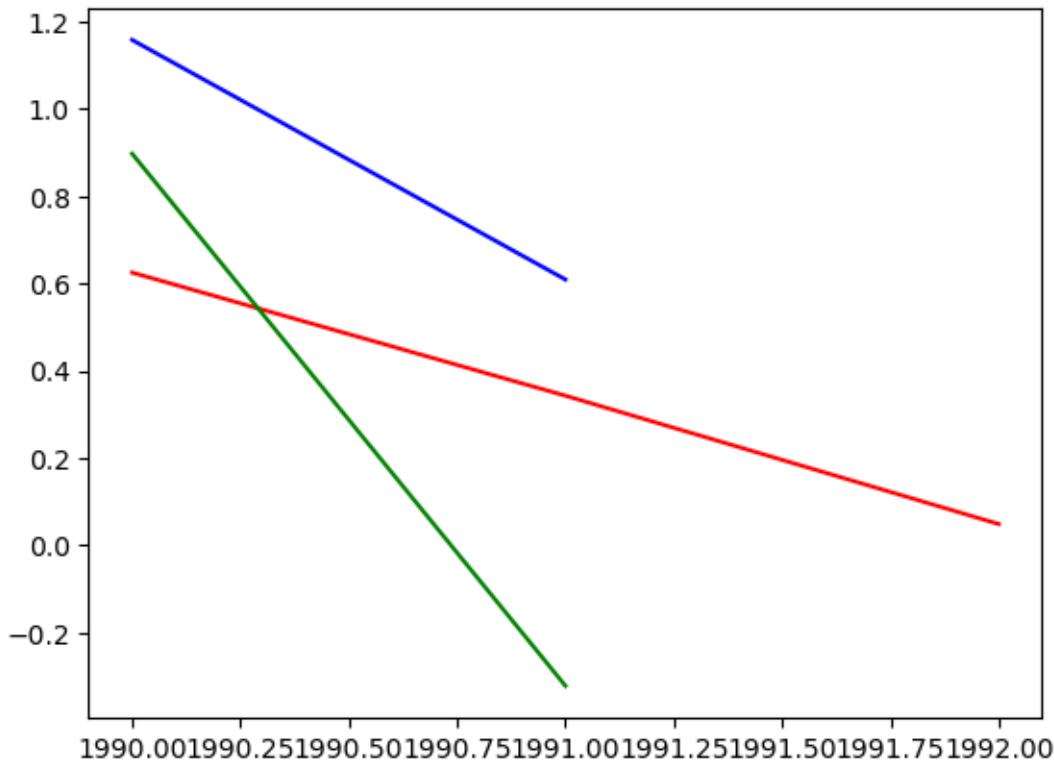
[5 rows x 31 columns]

```

```
[57]: # path 2 cont.

x_eth = df_cherry[df['area'] == 'Ethiopia PDR']['year']
y_eth = df_cherry[df['area'] == 'Ethiopia PDR']['average_temperature_c']
plt.plot(x_eth, y_eth, 'r')
x_ussr = df_cherry[df['area'] == 'USSR']['year']
y_ussr = df_cherry[df['area'] == 'USSR']['average_temperature_c']
plt.plot(x_ussr, y_ussr, 'b')
x_yugo = df_cherry[df['area'] == 'Yugoslav SFR']['year']
y_yugo = df_cherry[df['area'] == 'Yugoslav SFR']['average_temperature_c']
plt.plot(x_yugo, y_yugo, 'g')
plt.show()
```

```
/tmp/ipykernel_3250460/3954201872.py:2: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  x_eth = df_cherry[df['area'] == 'Ethiopia PDR']['year']
/tmp/ipykernel_3250460/3954201872.py:3: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  y_eth = df_cherry[df['area'] == 'Ethiopia PDR']['average_temperature_c']
/tmp/ipykernel_3250460/3954201872.py:5: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  x_ussr = df_cherry[df['area'] == 'USSR']['year']
/tmp/ipykernel_3250460/3954201872.py:6: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  y_ussr = df_cherry[df['area'] == 'USSR']['average_temperature_c']
/tmp/ipykernel_3250460/3954201872.py:8: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  x_yugo = df_cherry[df['area'] == 'Yugoslav SFR']['year']
/tmp/ipykernel_3250460/3954201872.py:9: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
  y_yugo = df_cherry[df['area'] == 'Yugoslav SFR']['average_temperature_c']
```



the data available for these countries is only from a few years

```
[58]: # observe year span of 1990-1991/1992
df_cherry
```

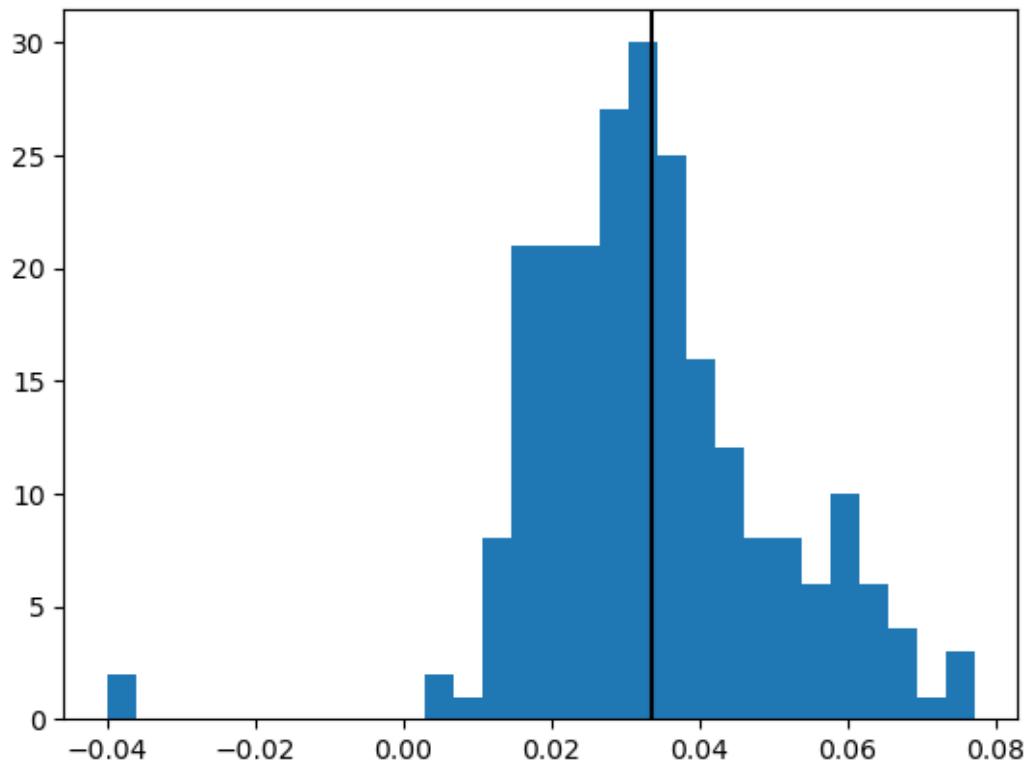
	area	year	savanna_fires	forest_fires	crop_residues	\
2152	Ethiopia PDR	1990	3851.3866	1509.6419	349.6276	
2153	Ethiopia PDR	1991	3851.3866	1509.6419	286.8652	
2154	Ethiopia PDR	1992	3851.3866	1509.6419	270.0245	
6684	USSR	1990	8405.2264	7262.4148	14854.7660	
6685	USSR	1991	8405.2264	7262.4148	11900.4386	
6901	Yugoslav SFR	1990	17.5124	27.8984	844.4150	
6902	Yugoslav SFR	1991	17.5124	27.8984	1047.9798	
			rice_cultivation	drained_organic_soils_co2	pesticides_manufacturing	\
2152			3391.507802	5513.0102	12.850182	
2153			3309.287814	5513.0102	12.098396	
2154			3586.696796	5513.0102	11.606165	
6684			4813.760000	131838.2352	1169.000000	
6685			4699.296000	131838.2352	1301.000000	
6901			69.619200	276.0991	40.000000	
6902			68.145300	276.0991	34.000000	

	food_transport	forestland	...	manure_management	\
2152	100.5070	-107.2750	...	1543.7229	
2153	98.8585	-107.2750	...	1554.3704	
2154	33.8336	-107.2750	...	1594.3863	
6684	32821.6716	-605722.9991	...	60407.5004	
6685	32500.3786	-605722.9991	...	59296.5706	
6901	1874.2445	-10843.0573	...	4584.6630	
6902	1727.6882	-10843.0573	...	4502.9731	
	fires_in_organic_soils	fires_in_humid_tropical_forests	...		\
2152	0.0			416.3368	
2153	0.0			416.3368	
2154	0.0			416.3368	
6684	0.0			0.0000	
6685	0.0			0.0000	
6901	0.0			0.0000	
6902	0.0			0.0000	
	on_farm_energy_use	rural_population	urban_population	...	\
2152	80.8401	44542203.0	6657624.0		
2153	82.1289	45925516.0	7022864.0		
2154	41.6932	47359180.0	7407170.0		
6684	248879.1769	99158922.0	188910774.0		
6685	246785.4967	99725999.0	189602194.0		
6901	836.4433	11929250.0	11445927.0		
6902	776.2953	11824752.0	11526722.0		
	total_population_male	total_population_female	total_emission	...	\
2152	24890395.0	25137639.0	65914.320714		
2153	25886400.0	26090250.0	61115.468490		
2154	26890139.0	27042568.0	61654.073430		
6684	136777703.0	153126995.0	524473.945541		
6685	137641632.0	153853429.0	520480.833615		
6901	10758632.0	11248553.0	28161.400700		
6902	10733004.0	11229798.0	24271.147832		
	average_temperature_c				
2152	0.625000				
2153	0.342667				
2154	0.048444				
6684	1.158250				
6685	0.609000				
6901	0.897500				
6902	-0.322417				

[7 rows x 31 columns]

```
[59]: # path 2
# slope distribution excluding outliers
slopes_filt = [i for i in slopes if i > -0.1]
plt.hist(slopes_filt,bins=30)
plt.axvline(x=np.mean(slopes_filt),color='black')
```

```
[59]: <matplotlib.lines.Line2D at 0x7f8167a5e780>
```



using regression analysis,

average temperature is rising in most areas after we filter out areas with small year span

what's next? we have many other dimensions to understand and formulate hypotheses with

```
[60]: df.head()
```

```
[60]:      area  year  savanna_fires  forest_fires  crop_residues \
0  Afghanistan  1990        14.7237       0.0557     205.6077
1  Afghanistan  1991        14.7237       0.0557     209.4971
2  Afghanistan  1992        14.7237       0.0557     196.5341
3  Afghanistan  1993        14.7237       0.0557     230.8175
4  Afghanistan  1994        14.7237       0.0557     242.0494
```

```

rice_cultivation drained_organic_soils_co2 pesticides_manufacturing \
0           686.00                  0.0          11.807483
1           678.16                  0.0          11.712073
2           686.00                  0.0          11.712073
3           686.00                  0.0          11.712073
4           705.60                  0.0          11.712073

food_transport forestland ... manure_management fires_in_organic_soils \
0           63.1152   -2388.803 ...      319.1763          0.0
1           61.2125   -2388.803 ...      342.3079          0.0
2           53.3170   -2388.803 ...      349.1224          0.0
3           54.3617   -2388.803 ...      352.2947          0.0
4           53.9874   -2388.803 ...      367.6784          0.0

fires_in_humid_tropical_forests on_farm_energy_use rural_population \
0                   0.0             NaN        9655167.0
1                   0.0             NaN       10230490.0
2                   0.0             NaN       10995568.0
3                   0.0             NaN       11858090.0
4                   0.0             NaN       12690115.0

urban_population total_population_male total_population_female \
0           2593947.0        5348387.0        5346409.0
1           2763167.0        5372959.0        5372208.0
2           2985663.0        6028494.0        6028939.0
3           3237009.0        7003641.0        7000119.0
4           3482604.0        7733458.0        7722096.0

total_emission average_temperature_c
0           2198.963539        0.536167
1           2323.876629        0.020667
2           2356.304229       -0.259583
3           2368.470529        0.101917
4           2500.768729        0.372250

```

[5 rows x 31 columns]

regression experiment showed the year span varies for diff areas.

these outliers may confound our analysis

what is the distribution of year span for each area?

## 1.9 Exercise 1

- goal: make a single figure to visualize and understand the year spans in the dataset
- time: 5m
- hint: you can use either 1D or 2D input for your plot

## 1.10 Population equilibrium

questions: - are male/female populations in equilibrium? - what is the rate of change of urban and rural populations?

```
[61]: df.head()
```

```
[61]:      area  year  savanna_fires  forest_fires  crop_residues \
0  Afghanistan  1990       14.7237      0.0557    205.6077
1  Afghanistan  1991       14.7237      0.0557    209.4971
2  Afghanistan  1992       14.7237      0.0557    196.5341
3  Afghanistan  1993       14.7237      0.0557    230.8175
4  Afghanistan  1994       14.7237      0.0557    242.0494

      rice_cultivation  drained_organic_soils_co2  pesticides_manufacturing \
0            686.00                  0.0          11.807483
1            678.16                  0.0          11.712073
2            686.00                  0.0          11.712073
3            686.00                  0.0          11.712073
4            705.60                  0.0          11.712073

      food_transport  forestland  ...  manure_management  fires_in_organic_soils \
0        63.1152   -2388.803  ...        319.1763          0.0
1        61.2125   -2388.803  ...        342.3079          0.0
2        53.3170   -2388.803  ...        349.1224          0.0
3        54.3617   -2388.803  ...        352.2947          0.0
4        53.9874   -2388.803  ...        367.6784          0.0

      fires_in_humid_tropical_forests  on_farm_energy_use  rural_population \
0                      0.0                NaN        9655167.0
1                      0.0                NaN        10230490.0
2                      0.0                NaN        10995568.0
3                      0.0                NaN        11858090.0
4                      0.0                NaN        12690115.0

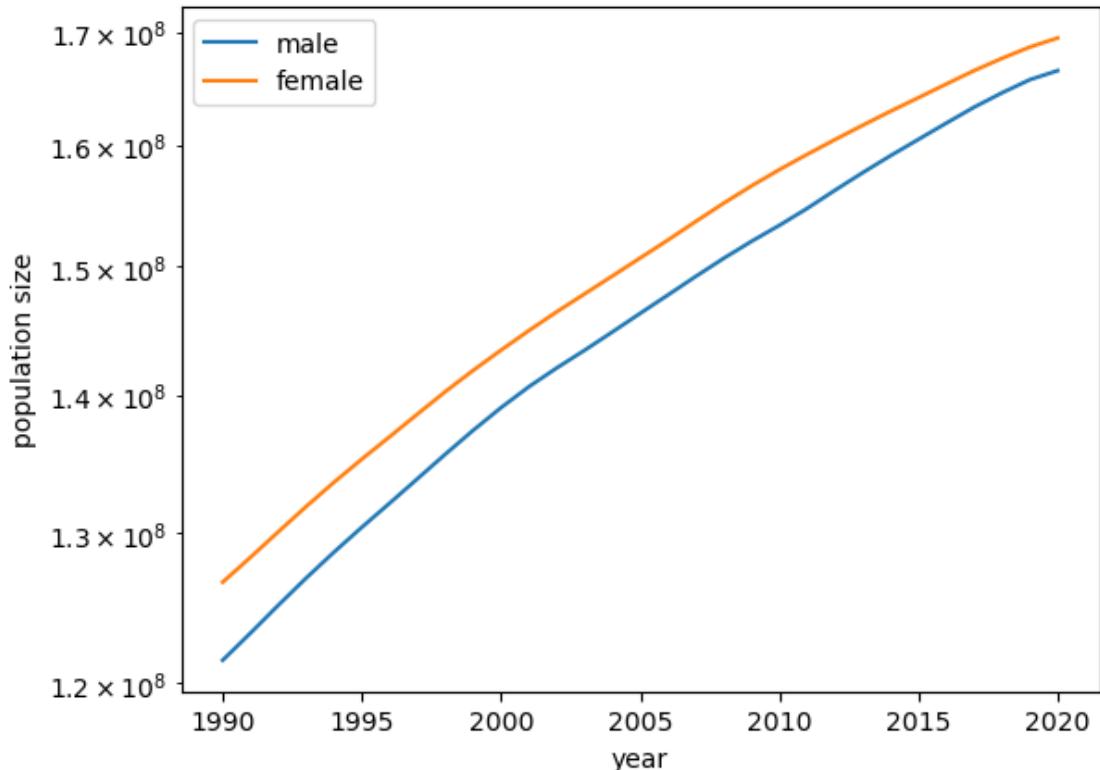
      urban_population  total_population_male  total_population_female \
0        2593947.0           5348387.0        5346409.0
1        2763167.0           5372959.0        5372208.0
2        2985663.0           6028494.0        6028939.0
3        3237009.0           7003641.0        7000119.0
4        3482604.0           7733458.0        7722096.0

      total_emission  average_temperature_c
0      2198.963539        0.536167
1      2323.876629        0.020667
2      2356.304229       -0.259583
3      2368.470529        0.101917
4      2500.768729        0.372250
```

[5 rows x 31 columns]

```
[62]: # plot male and female populations over time in us
x = df_usa['year']
y_male = df_usa['total_population_male']
y_female = df_usa['total_population_female']
plt.plot(x,y_male,label='male')
plt.plot(x,y_female,label='female')
plt.yscale("log",base=10)
plt.xlabel('year')
plt.ylabel('population size')
plt.legend()
```

[62]: <matplotlib.legend.Legend at 0x7f81727c21b0>

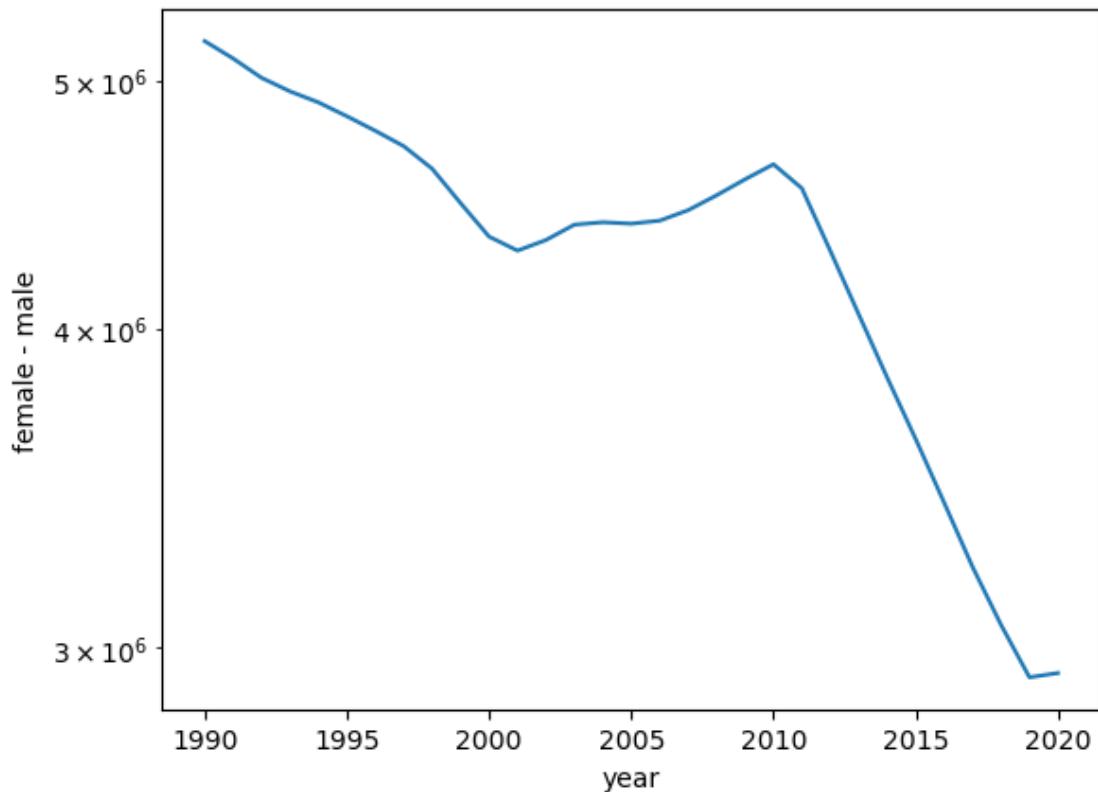


more females than males, but what is the gap here? hard to tell by this plot

```
[63]: # plot difference of males and females
y_diff = y_female - y_male
plt.plot(x,y_diff)
plt.yscale("log",base=10)
```

```
plt.ylabel('female - male')
plt.xlabel('year')
```

```
[63]: Text(0.5, 0, 'year')
```



- 1990: 5M more females than males
- 2020: only 3M more females than males
- the difference (females - males) is decreasing over time

## 1.11 Exercise 2

- goal: quantify the balance and rate of change between rural and urban populations in the USA
  - produce two line plots
1. Urban and rural population sizes over time
  2. The difference urban and rural populations over time

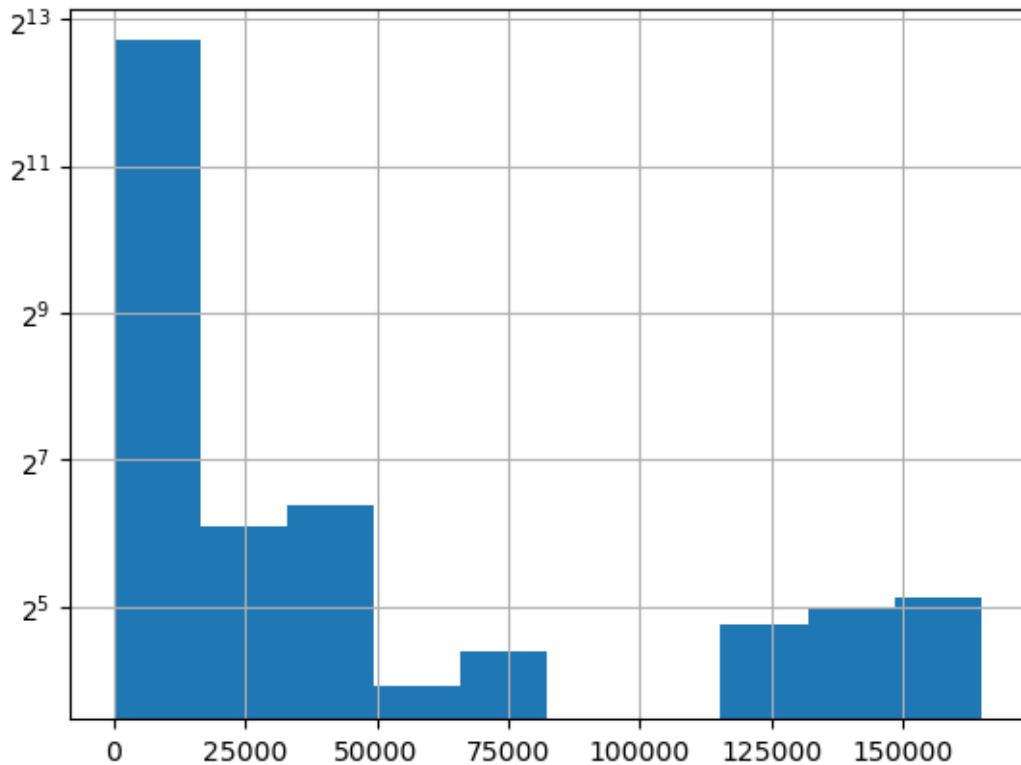
Write a 1 sentence hypothesis based on your plots

## 1.12 Rice cultivation dominated by few areas

the majority of human wealth and power is held by a small fraction of people

what about rice?

```
[64]: df['rice_cultivation'].hist()  
plt.yscale('log',base=2)
```



most areas produce a small amount of rice while relatively few produce a lot of rice

**heavy-tailed** distributions are infamous in social sciences and empirical network dataset

in fact, many non-negative count data sets (natural numbers  $\mathbb{N}_0$ ) are heavy-tailed

let's quantify how important the areas are on the right-hand tail of distribution

The **Gini coefficient** and **Lorenz curve** will quantify the degree of inequality in a distribution.

Lorenz curve

sort  $x_1, \dots, x_n$  from least to most “wealthy” - x-axis: cumulative fraction of the population - y-axis: cumulative fraction of total wealth

Gini coefficient -  $x \in \mathbb{R}^n$ : = data -  $\bar{x}$ : = empirical mean

$$G = \frac{\sum_{i=1}^n \sum_{j=1}^n |x_i - x_j|}{2n^2 \bar{x}}$$

- numerator is sum of pairwise differences
- divide by 2 bc we double count
- divide by  $n^2$  to average

- I feel like we should use  $\binom{n}{2}$  instead ...
- $G = \frac{\sum_{i=1}^n \sum_{j=1}^n |x_i - x_j|}{2n(n-1)\bar{x}}$
- divide by  $\bar{x}$  s.t.  $Range(G) \in [0, 1]$  (normalization)
- interpretation
  - $G = 0 \rightarrow$  wealth is equally distributed (one for all)
  - $G = 1 \rightarrow$  maximal inequality (all for one)

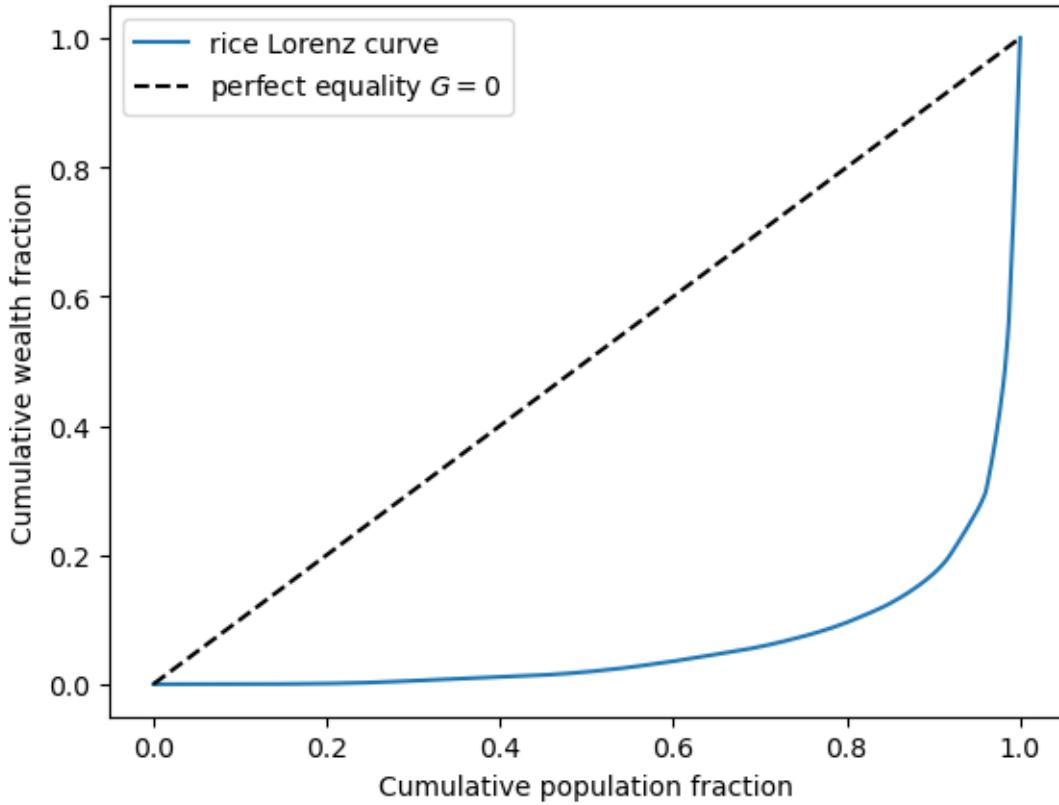
```
[65]: def gini(x, avg='nchoose2'):
    # there are more efficient ways to compute this
    # but the form used above is most interpretable
    n = len(x)
    num = 0
    mean = np.mean(x)
    for i in range(n):
        for j in range(n):
            num += np.abs(x[i] - x[j])
    if avg == 'nchoose2':
        denom = 2*n*(n-1) * mean
    elif avg == 'quad':
        denom = 2*(n**2)*mean
    else:
        raise ValueError('use "nchoose2" or "quad" for denominator scale')
    return num / denom
# takes about 1 min
gini(df['rice_cultivation'].values)
```

[65]: 0.8707302318270593

```
[66]: # plot lorenz curve
x = np.sort(df['rice_cultivation'].values)
n = len(x)
# cumulative sum vector
cs = np.cumsum(x)
# add 0,0
cs = np.insert(cs, 0, 0)
# normalize cs to [0,1]
cs = cs / cs[-1]
# space our population along the line segment [0,1]
pop = np.linspace(0,1,n+1) # +1 because we added 0,0
# plot
# lorenz curve
plt.plot(pop, cs, label='rice Lorenz curve')
# perfect equality line (G=0)
plt.plot([0,1], [0,1], label="perfect equality $G=0$", color = 'black', linestyle='--')
plt.xlabel('Cumulative population fraction')
plt.ylabel('Cumulative wealth fraction')
```

```
plt.legend()
```

[66]: <matplotlib.legend.Legend at 0x7f81678aee70>



interpretation of rice Lorenz curve - 20% of areas produce \$ 90y 0.1\$ of Lorenz curve)

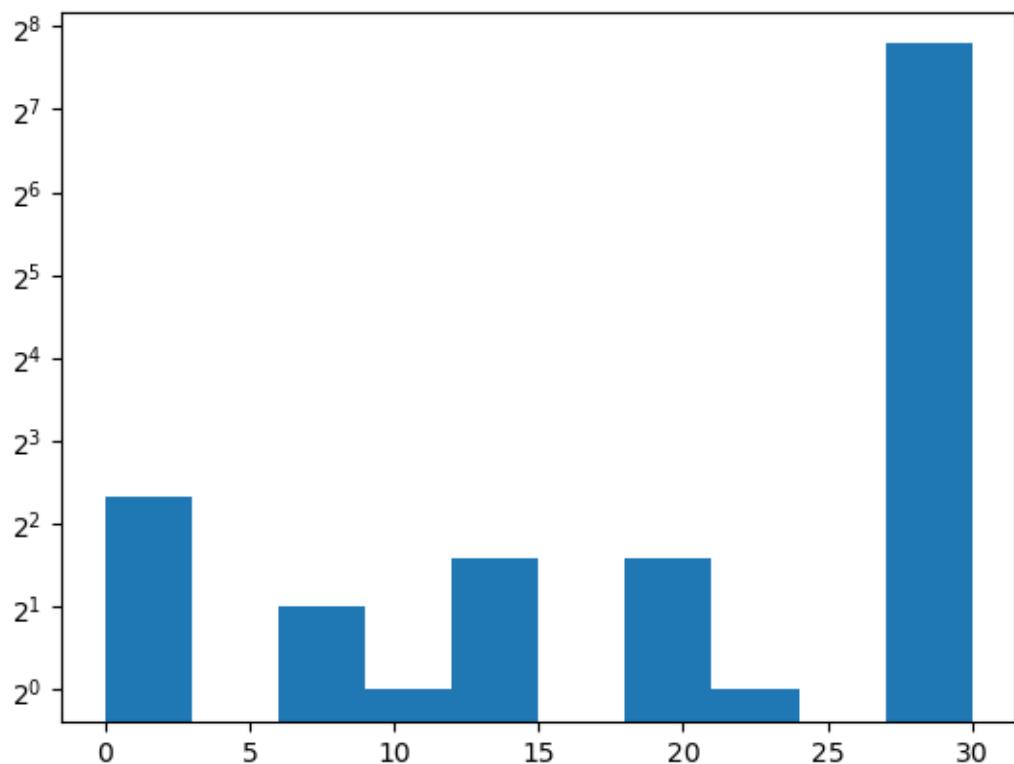
### 1.13 Exercise 3

- goal: compute Gini coef. for simulated data
- approach:
  - for zipf and powerlaw distributions (hint: use np.random)
    - \* sample 1000 values
    - \* compute Gini coef. and plot Lorenz curve
- time: 10m

## 2 Exercise solutions

### 2.1 Exercise 1

```
[67]: # histogram solution
spans = []
for group, df_group in df.groupby('area'):
    spans.append(df_group['year'].max() - df_group['year'].min())
plt.hist(spans)
plt.yscale('log',base=2)
```



histogram collapses year data for each area into a single integer: span length.

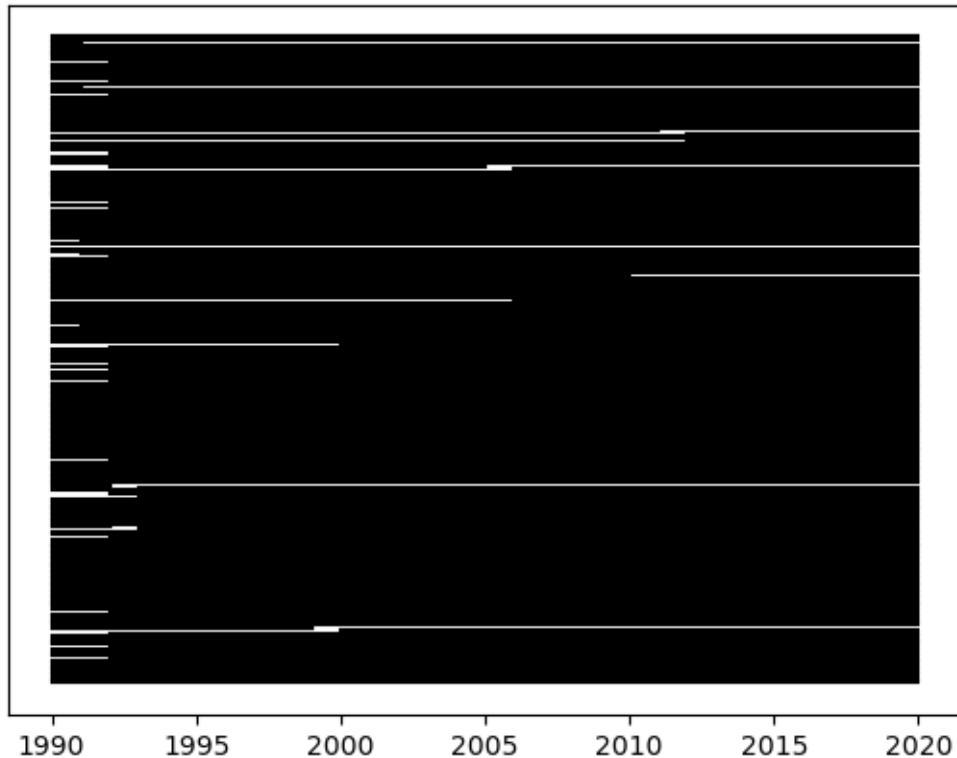
does not indicate which years are missing

```
[71]: # interval line solution
fig,ax = plt.subplots()
ticks = []
spacing = 1
i=-2
i_prev=0
for group, df_group in df.groupby('area'):
    i = i_prev + spacing
```

```

years = df_group['year'].values
ax.plot(years, np.repeat(i, len(years)), color='black')
ticks.append(group)
i_prev = i
# hide yticks
ax.yaxis.set_visible(False)
# set y axis label
ax.set_ylabel('Countries')
fig.show()

```



interval plot has fine grain resolution on which years are missing

## 2.2 Exercise 2

[72]: `df_usa.head()`

	area	year	savanna_fires	forest_fires	\
6591	United States of America	1990	1391.1481	1999.5617	
6592	United States of America	1991	1391.1481	1999.5617	
6593	United States of America	1992	1391.1481	1999.5617	
6594	United States of America	1993	1391.1481	1999.5617	
6595	United States of America	1994	1391.1481	1999.5617	

```

      crop_residues  rice_cultivation  drained_organic_soils_co2 \
6591      19193.1314          11195.520          50713.3546
6592      17574.1126          11029.312          50713.3546
6593      20964.8025          12421.402          50713.3546
6594      16761.5327          11235.602          50566.2744
6595      21370.6468          13151.110          50418.1314

      pesticides_manufacturing  food_transport  forestland ... \
6591                  13772.0          45410.2673 -520573.1593 ...
6592                  13139.0          44566.9408 -520573.1593 ...
6593                  14109.0          46198.5933 -520573.1593 ...
6594                  12948.0          45175.6740 -520573.1593 ...
6595                  14401.0          46739.0889 -520573.1593 ...

      manure_management  fires_in_organic_soils \
6591          47024.7024          0.0
6592          48206.0451          0.0
6593          48416.2439          0.0
6594          48426.0899          0.0
6595          49165.7227          0.0

      fires_in_humid_tropical_forests  on_farm_energy_use  rural_population \
6591                  9.8513          54454.7092          62373717.0
6592                  9.8513          54565.6091          61957131.0
6593                  9.8513          57469.9273          61539241.0
6594                  9.8513          54689.3180          61136396.0
6595                  9.8513          55539.8937          60759735.0

      urban_population  total_population_male  total_population_female \
6591          190156233.0          121451448.0          126632284.0
6592          193017688.0          123229931.0          128330258.0
6593          195915032.0          125081499.0          130093840.0
6594          198883790.0          126914111.0          131865642.0
6595          201981831.0          128685441.0          133588147.0

      total_emission  average_temperature_c
6591          463050.9394          0.733583
6592          473285.7816          0.706333
6593          486026.3425          0.253000
6594          484238.6197          0.153500
6595          509412.4984          0.470250

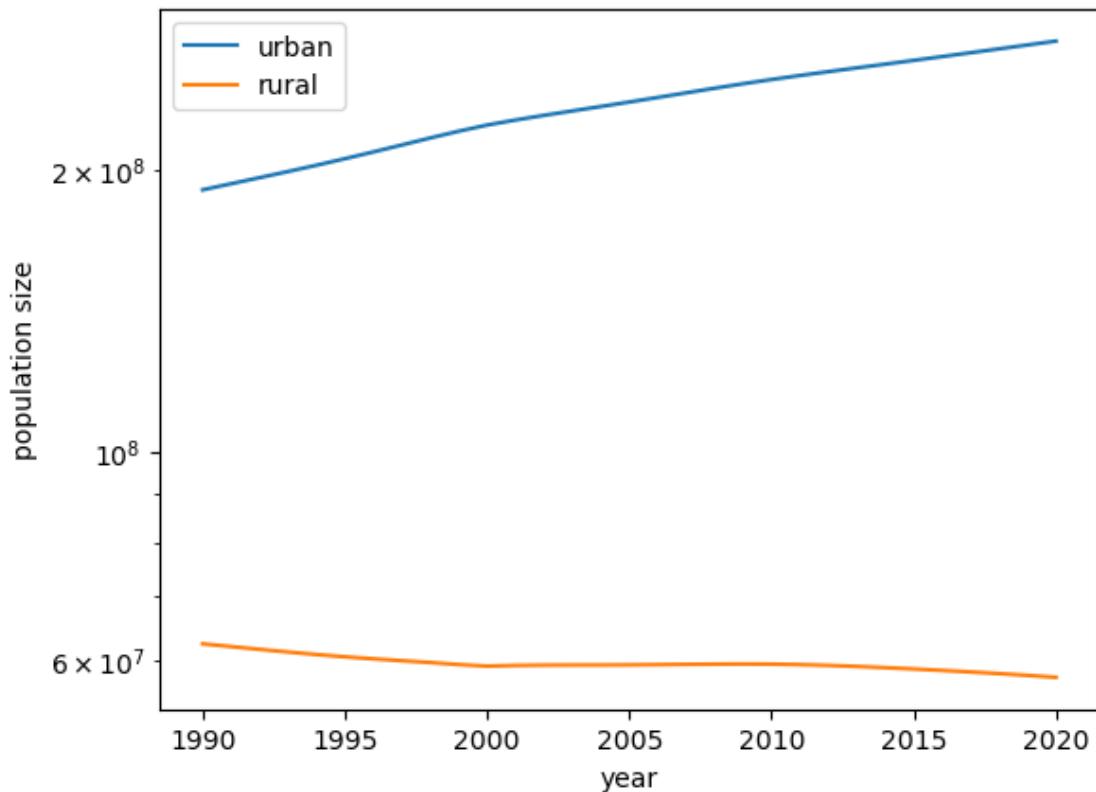
```

[5 rows x 31 columns]

```
[73]: # plot urban and rural populations over time in us
x = df_usa['year']
```

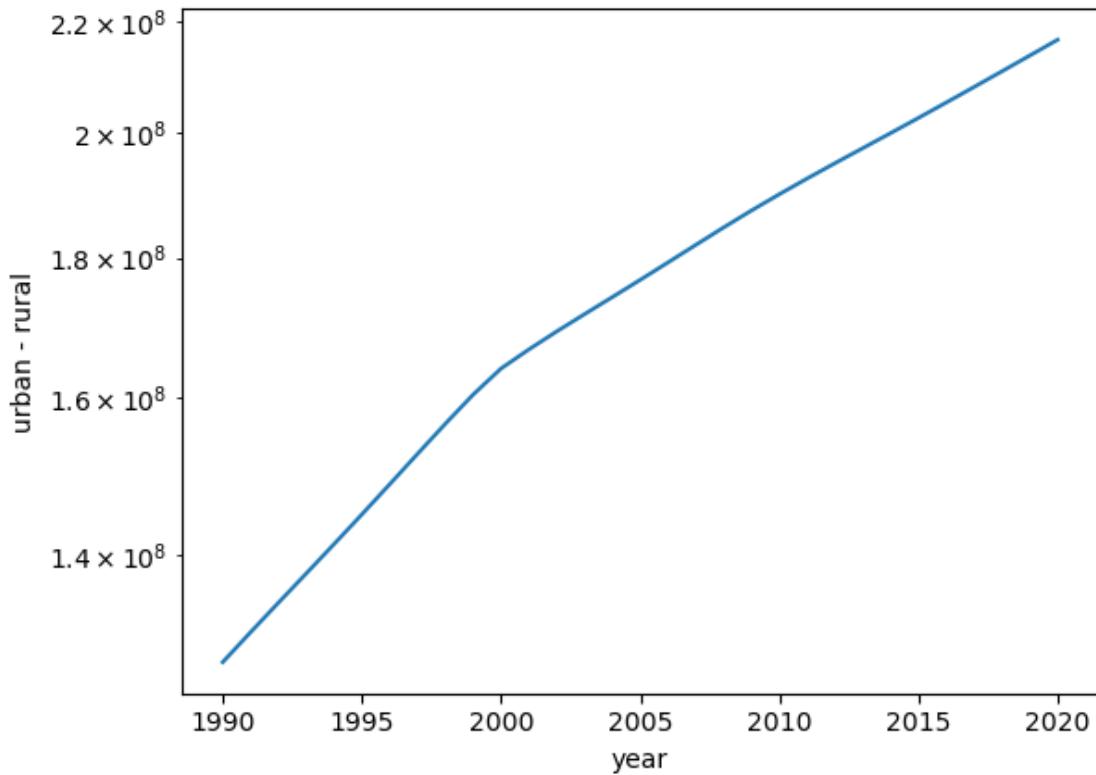
```
y_urban = df_usa['urban_population']
y_rural = df_usa['rural_population']
plt.plot(x,y_urban,label='urban')
plt.plot(x,y_rural,label='rural')
plt.yscale("log",base=10)
plt.xlabel('year')
plt.ylabel('population size')
plt.legend()
```

[73]: <matplotlib.legend.Legend at 0x7f8171481d00>



```
[74]: # plot difference of males and females
y_diff = y_urban - y_rural
plt.plot(x,y_diff)
plt.yscale("log",base=10)
plt.ylabel('urban - rural')
plt.xlabel('year')
```

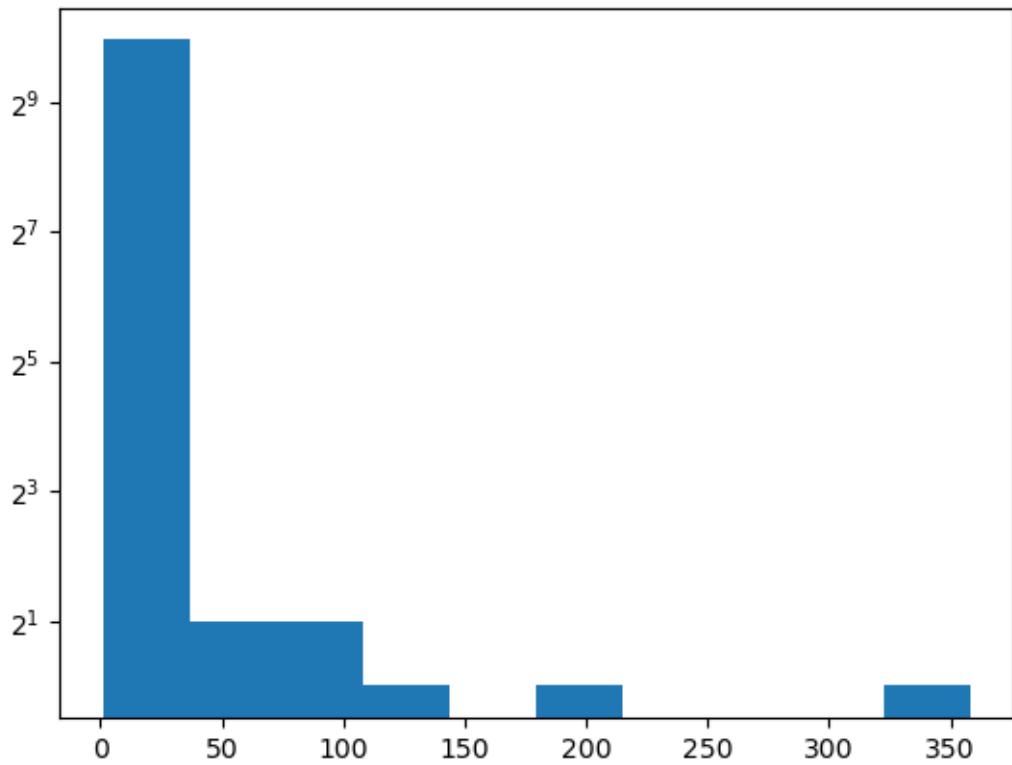
[74]: Text(0.5, 0, 'year')



hypothesis: in the last three decades of the USA, people have been migrating from rural to urban areas

### 2.3 Exercise 3

```
[75]: x_zipf = np.random.zipf(a=2,size=1000)
plt.hist(x_zipf)
plt.yscale("log",base=2)
```



```
[76]: # zipf Gini
gini_zipf = gini(x_zipf)
print("zipf gini: ", gini_zipf)
```

zipf gini: 0.6202736694539973

```
[77]: # zipf Lorenz
x_zipf_sort = np.sort(x_zipf)
n = len(x_zipf_sort)
# cumulative sum vector
cs = np.cumsum(x_zipf_sort)
# add 0,0
cs = np.insert(cs, 0, 0)
# normalize cs to [0,1]
cs = cs / cs[-1]
# space our population along the line segment [0,1]
pop = np.linspace(0,1,n+1) # +1 because we added 0,0
# plot
# lorenz curve
plt.plot(pop, cs, label='Zipf Lorenz curve')
# perfect equality line (G=0)
```

```

plt.plot([0,1], [0,1], label="perfect equality $G=0$", color = 'black',  

         linestyle='--')  

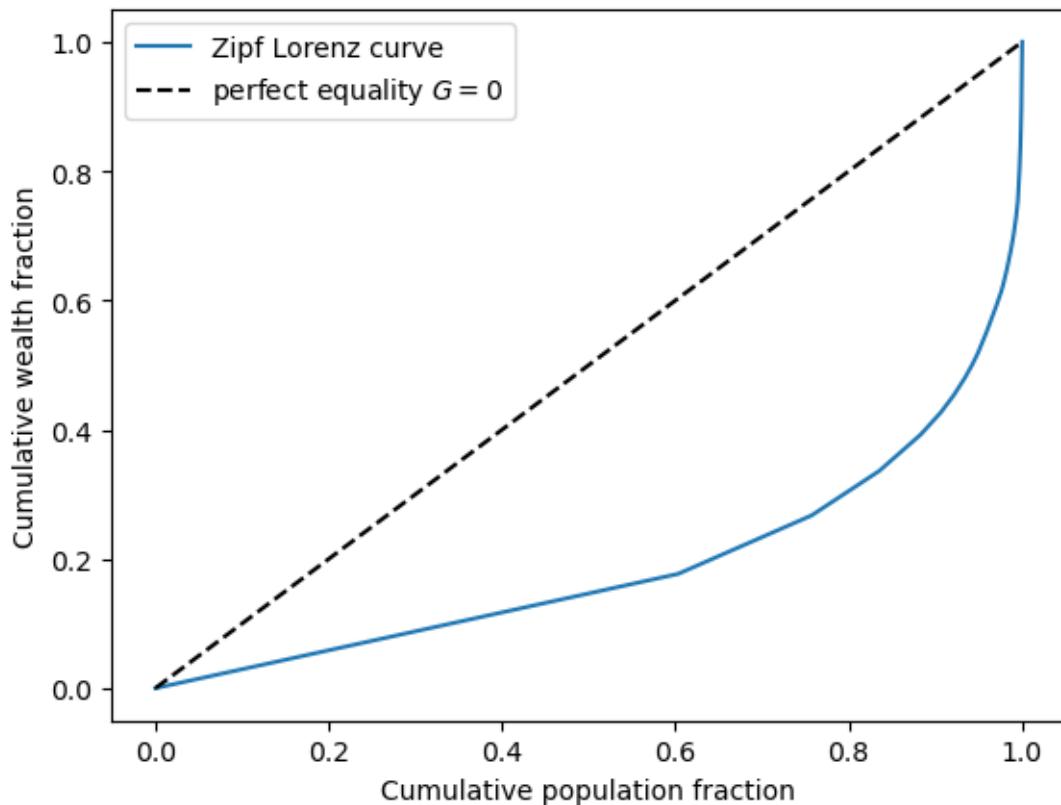
plt.xlabel('Cumulative population fraction')  

plt.ylabel('Cumulative wealth fraction')  

plt.legend()

```

[77]: <matplotlib.legend.Legend at 0x7f8171735430>



```

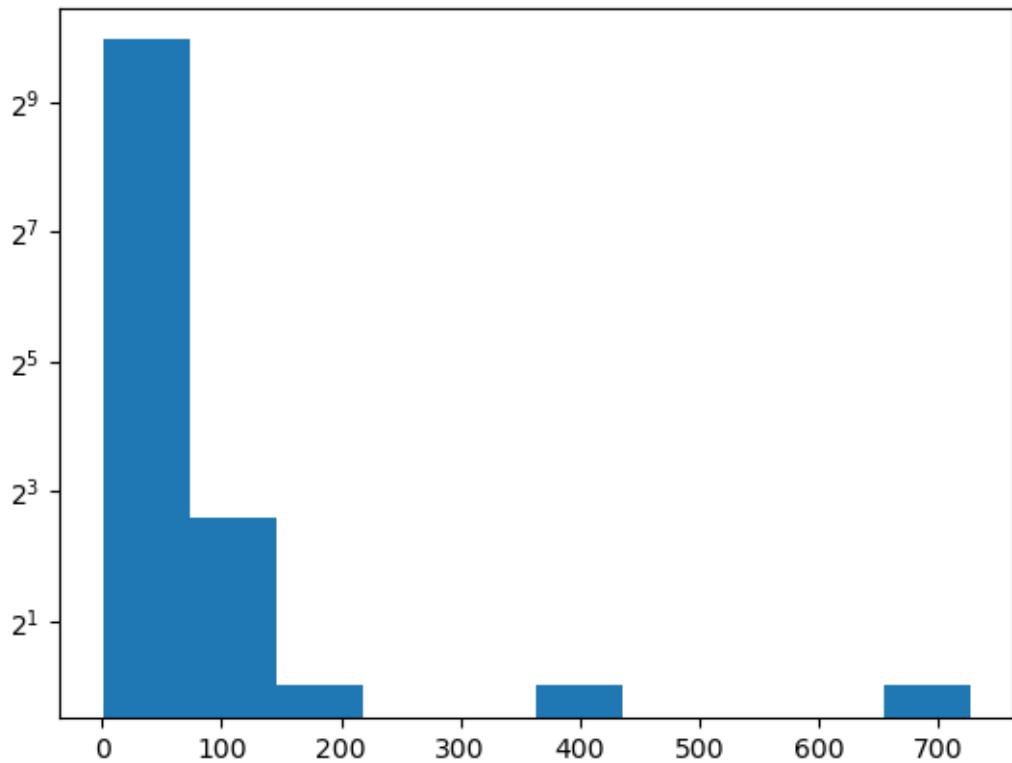
[78]: # zipf a = 2  

x_zipf = np.random.zipf(2,1000)  

plt.hist(x_zipf)  

plt.yscale("log",base=2)

```



```
[79]: # zipf Gini
gini_zipf = gini(x_zipf)
print("zipf gini: ", gini_zipf)
```

zipf gini: 0.739825818556948

```
[80]: # zipf Lorenz
x_zipf_sort = np.sort(x_zipf)
n = len(x_zipf_sort)
# cumulative sum vector
cs = np.cumsum(x_zipf_sort)
# add 0,0
cs = np.insert(cs, 0, 0)
# normalize cs to [0,1]
cs = cs / cs[-1]
# space our population along the line segment [0,1]
pop = np.linspace(0,1,n+1) # +1 because we added 0,0
# plot
# lorenz curve
plt.plot(pop, cs, label='Zipf Lorenz curve')
# perfect equality line (G=0)
```

```

plt.plot([0,1], [0,1], label="perfect equality $G=0$", color = 'black',  

         linestyle='--')  

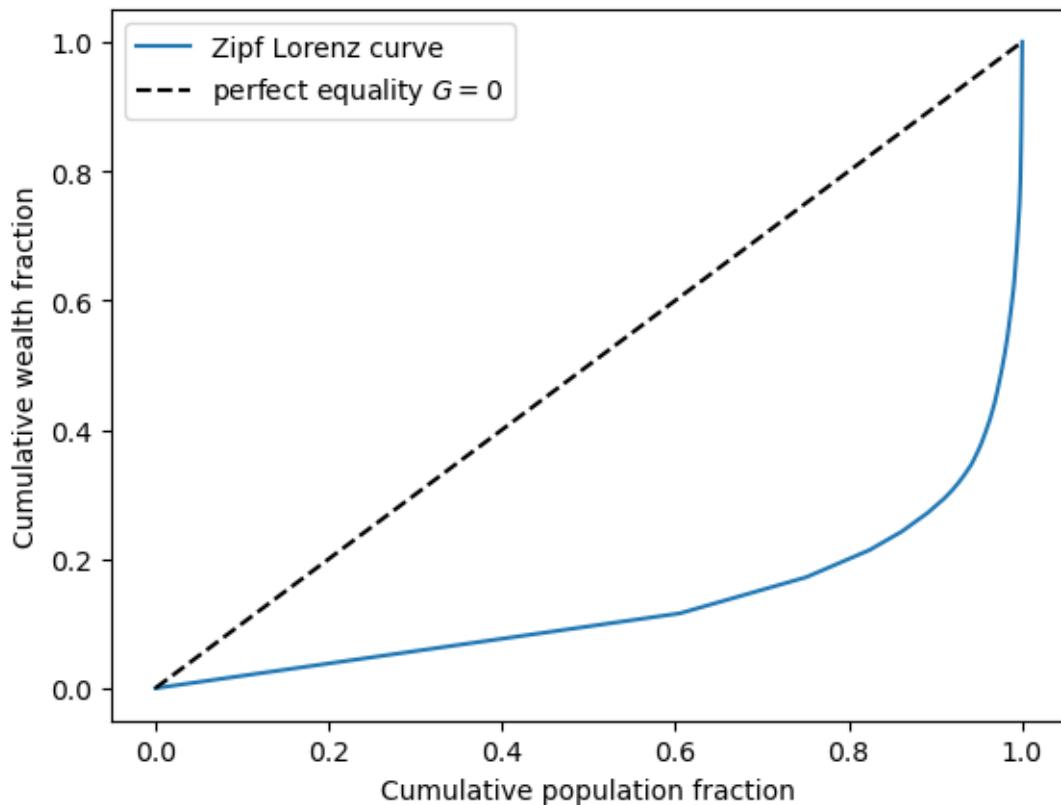
plt.xlabel('Cumulative population fraction')  

plt.ylabel('Cumulative wealth fraction')  

plt.legend()

```

[80]: <matplotlib.legend.Legend at 0x7f8167279640>



increasing the  $a$  parameter increased the Gini Coefficient and changed the shape of Lorenz curve.

### 3 Instructor notes

from command-line

jupyter server

copy link into web browser