



# Python Libraries


CSCI 6118 Software Engineering for Scientists  
Jake Krol 2025





# Introduction

- Position: Computer science Ph.D. student in Layer lab
- Research: population genomics




# Outline

- Purpose of python libraries
- Terminology
- Brief history
- NumPy
- Pandas
- Matplotlib
- Physical structure of packages



## With respect to science

The Scientific Method as an Ongoing Process




Can we  
cycle faster  
using  
libraries?

# Learning goals


- Why are **libraries good for science?**
- What do “**modules**”, “**packages**”, and “**libraries**” mean w.r.t Python?
- How/when use **NumPy, Pandas, Matplotlib**?
- What alternatives exist for **big data scenarios**?
- What is the **physical structure** of a python package?

# Purpose of libraries

- Formally, a library is a code for general purpose tasks
- $\text{utility}_{\text{lib}}(\text{trust}_{\text{lib}}, \text{time}, \text{necessity})$ 
  - $\text{trust}_{\text{lib}}$
  - time
  - necessity




The Scientific Method as an Ongoing Process



# Terminology

**Libraries** are composed of **packages** which are composed of **modules**



# History of Python packages

- Popular libraries are often “old” and have high *utility*
  - $\text{utility}_{\text{lib}}(\text{trust}_{\text{lib}}, \text{time}, \text{necessity})$



# NumPy

- Purpose
  - Init, view, modify, or write multi dimensional arrays
  - Random number generation



# NumPy

- Array initializations. Defining and viewing shape

```
l = list(range(0,12))
```

```
a = np.array(l) # vector
```

```
np.reshape(a, (3,4)) # matrix
```

```
a = np.array(l).reshape(4,3) # matrix
```

```
a = np.array(l).reshape(2,3,2) # tensor
```




Q: What's the difference?

```
np.reshape(a,(3,4), 'C')  
np.reshape(a,(3,4), 'F')
```

# NumPy

- Check shape
  - `np.shape(a)`
- Check # dimensions
  - `np.ndim(a)`
- Q: What is output shape of these two operations?
  - `n = len(a) # a is a vector`
  - **`np.dot(a,a)`**
  - **`np.outer(a,a)`**



# NumPy

- Elementwise operations

$a + 5$

$a - 5$

$a * 10$

$a / 2$

- NumPy is usually more efficient than base python approaches
  - Low-level C code performs the operations on elements without looping

```
a = np.arange(1000000) # 1M
```



```
t0 = time.time()
b1 = [i * 2 for i in a]
t1 = time.time()
print("List comprehension:", round(t1 - t0,2),
"seconds")
t0 = time.time()
b2 = a * 2
t1 = time.time()
print("NumPy vectorized:", round(t1 - t0,2),
"seconds")
List comprehension: 0.065 seconds
NumPy vectorized: 0.002 seconds
```

**Are 1M elements  
realistic?**



# NumPy

- Yes, 1M elements are *common* in scientific applications
- Ex: genomics
  - Say, we inspect 1/10,000th of human genome for variants
    - ~300k base pairs
  - Compare only 10 people
  - We have  $(3 \times 10^5 \text{ variants}) \times (10 \text{ people}) = 3 \times 10^6$



My research  
compares  
>1000 human  
genomes

Variant  
↓  
TAGCG ... AGTCT  
ATCGC ... TCAGA  
~3 **billion** base pairs in  
human reference genome

# NumPy

- How do non-trivial math quickly on arrays?
  - Numba



```
def f(x):  
    if x > 0:  
        y = np.log(x)  
    else:  
        y = x ** 2
```

# NumPy v Numba

\*Conditional logic  
must operate on  
scalars only

Q: how vectorize a user defined function?

...

```
x = np.linspace(-1000, 1000, 10_000_000) # data

def piecewise_numpy(x): # original function
    return np.where(x > 0, np.sqrt(x), -np.sqrt(-x))

start = time.time()

y_numpy = piecewise_numpy(x)

numpy_time = time.time() - start
```

...

```
@vectorize(["float64(float64)"], target='parallel')
# decorator

def piecewise_numba_vec(x): # numba function
    lambda x: np.sqrt(x) if x > 0 else -np.sqrt(-x)

piecewise_numba_vec(x[:10]) # compile func by
subset run

start = time.time()

y_numba_vec = piecewise_numba_vec(x)

numba_vec_time = time.time() - start
```




**Original:** 0.09s  
**Numba:** 0.01s

# Python v NumPy v Numba



- Python takes ~x100 longer than Numba
- NumPy takes ~x10 longer than Numba

The Scientific Method as an Ongoing Process



# NumPy

- Random number generation
- Applications
  - Monte carlo hypothesis testing
  - Simulations
- Simulations are a cheap, fast way to test experimental setups, hypotheses, etc.
  - **Expected data can test expected results**
  - Proactively setup pipelines before real data is collected



# NumPy

```
n_flips = 20
```

```
heads_emp = 15
```

```
n_permutations = 10000
```

```
# simulate 20 fair coins flips 10k times
```


```
simulated_heads = np.random.binomial(n_flips,  
0.5, n_permutations)
```

```
# compute one-sided p-value (alternative: more  
heads than expected)
```

```
num_null_extreme = np.sum(simulated_heads  
>= empirical_heads)
```

```
p_value = (num_null_extreme + 1) /  
(n_permutations + 1) # +1 for empirical
```

**One-sided p-value: 0.0225**



# Numpy

- Support for various distribution functions

?np.random ->

Univariate  
distributions

---

beta	Beta distribution over ``[0,1]``.
binomial	Binomial distribution.
chisquare	:math:`\chi^2` distribution.
exponential	Exponential distribution.
f	F (Fisher-Snedecor) distribution.
gamma	Gamma distribution.
geometric	Geometric distribution.
gumbel	Gumbel distribution.
hypergeometric	Hypergeometric distribution.
laplace	Laplace distribution.
logistic	Logistic distribution.
lognormal	Log-normal distribution.
logseries	Logarithmic series distribution.
negative_binomial	Negative binomial distribution.
noncentral_chisquare	Non-central chi-square distribution.
noncentral_f	Non-central F distribution.
normal	Normal / Gaussian distribution.
pareto	Pareto distribution.
poisson	Poisson distribution.
power	Power distribution.
rayleigh	Rayleigh distribution.
triangular	Triangular distribution.
uniform	Uniform distribution.
vonmises	Von Mises circular distribution.
wald	Wald (inverse Gaussian) distribution.
weibull	Weibull distribution.
zipf	Zipf's distribution over ranked data.



# NumPy

- Summary
  - For “box” data
  - Fast numeric computation
  - Simulation from various distributions
- Advanced
  - Vectorized functions for large datasets (numba or built-ins)
  - Specialized compression formats for faster I/O
    - Look into `np.savez(file, arrays)`
- Limitations
  - Best for *numeric* data
  - Pandas is usually better for tables with mixed data types (strings + numbers)





# Pandas

- Purpose
  - Read, init, view, modify, or write tabular data
- How Pandas differ from NumPy?
  - **Optimal for 2D** data only
  - Better at “**nuancing**” which values undergo an **operation**
    - e.g., `groupby()`
- When applicable, Pandas will use NumPy for numeric operations

# Pandas

- Basics

```
df = pd.DataFrame( {"x": [1, 2], "y": [3,4]} ) # init
```

```
df.to_csv("example.tsv", index=False, sep='\t') # write
```

```
df = pd.read_csv("example.tsv", sep='\t') # read
```

```
df['z'] = df.apply(axis=1, func = lambda row: row['x'] + row['y']) # row-wise operation
```

```
df['u'] = df['x'].apply(func=lambda i: i**2) # column-wise operation
```

- So far, similar to NumPy.

# Pandas

- LOCating values

```
df = pd.DataFrame( {"x": [1, 2], "y": [3,4]} ) # init
```

```
df.iloc[1,0] # numeric indexing (0-based)
```

Q: which value is returned?


```
df.loc[0,'x'] # name indexing (by default, row indices are already integers)
```

```
df['x'] # by default, name indexing searches the columns for matching
```




```
df.columns # column name extraction
```

# Pandas

- Common uses
  - **Key**-dependent operations
    - Groupby(<**key**>)
    - pd.merge(x,y,on=<**key**>)
    - <**key**> is a column name
  - Again, Pandas is great for “nuancing” which data subsets operations apply to





Group  
then sum



# Pandas

- In the context of Pandas, **keys** are tags that categorize items (usually rows).



# Pandas

- Groupby
  - 1. Collect all rows with shared key
  - 2. Perform a *summary* operation on groups

```
df = pd.DataFrame({'key': ['a','a','b','b'], 'val': [-1,1,2,4]})
```

```
grouped = df.groupby('key')
```


```
grouped.sum()
```

```
# outputs a len(unique_keys) x 2 table
```

```
# key      value
```

```
# a      0
```

```
# b      6
```




What other group summary funcs are built-in?

```
print(';'.join(dir(grouped)))
```

agg;aggregate;all;any;apply;bfill;boxplot;corr;corrwith;count;  
cov;cumcount;cummax;cummin;cumprod;cumsum;describe;  
diff;dtypes;ewm;expanding;ffill;fillna;filter;first;get\_group;grou  
ps;head;hist;idxmax;idxmin;indices;key;last;**max;mean;med  
ian;min;ndim;ngr**oup;ngrou**p;nth;nunique;ohlc;pct\_change;  
pipe;plot;prod;quantile;rank;resample;rolling;sample;sem;shi  
ft;size;skew;std;sum;tail;take;transform;value\_counts;var**

# Pandas

- Join
  - Combine  $\geq 2$  datasets by a shared key
  - Types
    - Inner: intersection
    - Left: A and intersection
    - Right: B and intersection
    - Outer: All keys



# Pandas

```
A = pd.DataFrame({  
    "key": ["a", "a", "b", "b"],  
    "val_a": [-1, 1, 2, 4]  
})  
  
B = pd.DataFrame({  
    "key": ["a", "b", "c"],  
    "val_b": ['x', 'y', 'z']  
})  
  
pd.merge(A,B,on='key',how='left')
```

Q: What shape will the output be?

# Pandas

```
A = pd.DataFrame({  
    "key": ["a", "a", "b", "b"],  
    "val_a": [-1, 1, 2, 4]  
})  
  
B = pd.DataFrame({  
    "key": ["a", "b", "c"],  
    "val_b": ['x', 'y', 'z']  
})  
  
pd.merge(A,B, on='key', how='left')
```

```
# 4 x 3  
key  val_a  val_b  
a    -1     x  
a     1     x  
b     2     y  
b     4     y
```

Q: what shape would an outer join be?


A: 5 x 3

# Pandas

- Limitation
  - Row-wise operations can be very slow
    - Solution: swifter
  - I/O is slow for large data and memory usage is high
    - Alternative: duckdb



## The Scientific Method as an Ongoing Process



# Pandas

- Use **swifter** library for **really easy rowwise parallelization**

```
# used 50 cores for experiment
N = 10_000_000
df = pd.DataFrame({
    'a': np.random.rand(N),
    'b': np.random.rand(N),
    'c': np.random.rand(N)
})
def rowf(row):
    return math.exp(row.a) \
        * math.sin(row.b) \
        + math.log1p(row.c**2)
```

```
result_pandas = df.apply(rowf, axis=1)
result_swifter = df.swifter.apply(rowf,
axis=1)
```




**Pandas:** 97.39 s  
**Swifter:** 16.27 s

# Pandas

- Summary
  - 2D data
  - Better non-numeric support than NumPy
    - Keys and indexing enable ...
      - “Nuanced” operation scope
      - Relate  $\geq 2$  datasets
- Pandas has some built-in plotting support
  - Ex: `df['x'].hist()`
- However, matplotlib code is required to control format

# Matplotlib




- Purpose
  - Data visualization
- For science communication,
  - Plots should be easy to interpret
  - Science is complex, and the goal of communication is to convey key takeaways to our audience




# Matplotlib

- 1-dimensional input data
  - Histogram
  - Boxplot
  - Violin

Q: Which should you choose?



# Matplotlib



## Use cases

- Distribution analysis for ~<=3 empirical vectors
- Median and IQR comparison for many empirical vectors
- Distribution analysis for ~>3 empirical vectors
- My least fav.

# Matplotlib

- 1-dimensional plotting code

```
import matplotlib.pyplot as plt; import numpy as np
```


```
data = np.random.randn(1000)
```

```
plt.hist(data,bins=30,color='skyblue',edgecolor='black')
```

```
plt.title('Histogram')
```

```
plt.show() # for interactive sessions
```



```
plt.savefig('hist.png') # save to png file
```



# Matplotlib



- 2-dimensional input data
  - Scatter
  - Line

Q: what would line plot of left look like?



# Matplotlib

- **Line** plots make sense for **one-to-one** mappings (proper functions)
- **Scatter** plots are good for **empirical data**



# Matplotlib

- Scatter and line plotting code

```
### scatter
```


```
# standard normal sampling
```

```
x = np.random.randn(100)
```

```
y = np.random.randn(100)
```

```
plt.scatter(x,y)
```

```
plt.show()
```



```
### line
```

```
# 100 values eq spaced from [0,10]
```

```
x = np.linspace(0,10,100)
```



```
y = np.sin(x)
```

```
plt.plot(x,y)
```

```
plt.show()
```

# Matplotlib

- 3-dimensional input data
  - Heatmap



# Matplotlib

- Heatmap plotting code


```
# 5x5 array sampled from [0,1) unif dist
```

```
data = np.random.rand(5,5)
```

```
# heatmaps are equivalently images
```

```
plt.imshow(data,cmap='viridis')
```

```
plt.colorbar(label='Value')
```



# Matplotlib

- Title, axes, and formatting

`plt.title()`

`plt.xlabel()`

`plt.ylabel()`

*# set size (inches) upon init*

`plt.figure(figsize=(8,6))`

*# get axes object of plot*

`ax = plt.gca()`





*# rm top boundary line*

`ax.spines['top'].set_visible(False)`

# Matplotlib

- Object oriented approach has finer control, but diff syntax
- Allows for subplots


```
fig, axes = plt.subplots(2, 2, figsize=(10, 4))
axes[0,0].hist(np.random.randn(1000))
axes[0,1].scatter(np.random.randn(100),np.random.randn(100))
axes[1,0].plot(np.linspace(0,10,100), np.sin(np.linspace(0,10,100)))
axes[1,1].imshow(np.random.rand(5,5), cmap='viridis')
```




# Matplotlib v Plotly

- Limitation
  - Static outputs
- Alternative plotly
  - \***Caution:** interactive plots are *good for doing science, but not great for communicating science*
    - Do you expect a reviewer to zoom/pan to the correct location?
      - They may not even open the file
    - Consider using gifs/videos if using interactive plots in presentations


# Matplotlib v Plotly



## The Scientific Method as an Ongoing Process






$\text{utility}_{\text{lib}}(\text{trust}_{\text{lib}}, \text{time}, \text{necessity})$



# Matplotlib

- Summary
  - The goto python visualization library
  - Examples for 1/2/3-dimensional input data
  - Best for static plots
    - Plotly is an interactive alternative
- Remember
  - Plots should communicate science efficiently; beware complex plots



# Beyond NumPy Pandas Matplotlib

- Scipy: stats and optimization
- Scikit-learn: machine learning
- Pytorch: deep learning
- Networkx: network analysis
- Pyvis: interactive network visualization

*Which libraries will emerge next?*



# Physical structure of package

- Scenario: You need to write/edit a python package for your science project
- Approach
  - 1. Install an editable version
  - 2. Implement changes
  - 3. Deploy
- Free deployment at <https://pypi.org/>




# Physical structure of package

- Editing pandas library

*git clone*

<https://github.com/pandas-dev/pandas.git>



```
jake[0]:~/ghub/layer-lab/teaching/pandas main • tree -L 1 .  
*  
|   asv_bench  
|   AUTHORS.md  
|   ci  
|   CITATION.cff  
|   codecov.yml  
|   doc  
|   Dockerfile  
|   environment.yml  
|   generate_pxi.py  
|   generate_version.py  
|   gitpod  
|   LICENSE  
|   LICENSES  
|   MANIFEST.in  
|   meson.build  
|   pandas  
|   pyproject.toml  
|   pyright_reportGeneralTypeIssues.json  
|   README.md  
|   requirements-dev.txt  
|   scripts  
|   setup.py  
|   tooling  
|   typings  
|   web
```


# Physical structure of package

- **pyproject.toml** OR **setup.py** have pkg metadata

```
[project]
name = 'pandas'
dynamic = [
    'version'
]
description = 'Powerful data structures for data analysis, time series, and statist
ics'
readme = 'README.md'
authors = [
    { name = 'The Pandas Development Team', email='pandas-dev@python.org' },
]
license = {file = 'LICENSE'}
requires-python = '>=3.11'
dependencies = [
    "numpy>=1.26.0",
    "python-dateutil>=2.8.2",
    "tzdata>=2023.3"
```

# Physical structure of package

- Core
  - Key functions and classes
- `__init__.py`
  - Declares as package



```
jake[0]:~/ghub/layer-lab/teaching/pandas/pandas main ✦ tree -L 1
.
├── api
├── arrays
├── compat
├── _config
├── conftest.py
├── core
├── errors
└── __init__.py
├── io
├── _libs
├── meson.build
├── plotting
├── _testing
├── testing.py
├── tests
├── tseries
├── _typing.py
├── util
└── _version.py
```

# Physical structure of package

- What's inside core package?
  - [frame.py](#), [series.py](#), ...

```
(pandas_venv) jake[0]:~/ghub/layer-lab/teaching/pandas/pandas/core main $ tree -L 1
├── accessor.py
├── algorithms.py
├── api.py
├── apply.py
├── array_algos
├── arraylike.py
├── arrays
├── base.py
├── col.py
├── common.py
├── computation
├── config_init.py
├── construction.py
├── dtypes
├── flags.py
├── frame.py
├── generic.py
├── groupby
├── indexers
├── indexes
├── indexing.py
├── __init__.py
├── interchange
├── internals
├── methods
├── missing.py
├── nanops.py
├── _numba
├── ops
├── resample.py
├── reshape
├── roperator.py
├── sample.py
├── series.py
├── shared_docs.py
├── sorting.py
├── sparse
├── strings
├── tools
└── util
└── window
```

# Physical structure of package

- Adding a print statement to shape attribute

```
@property
def shape(self) -> tuple[int, int]:
    """
    Return a tuple representing the dimensionality of the DataFrame

    Unlike the `len()` method, which only returns the number of rows,
    provides both row and column counts, making it a more informative
    understanding dataset size.

    See Also
    -----
    numpy.ndarray.shape : Tuple of array dimensions.

    Examples
    -----
    >>> df = pd.DataFrame({"col1": [1, 2], "col2": [3, 4]})
    >>> df.shape
    (2, 2)



    >>> df = pd.DataFrame({"col1": [1, 2], "col2": [3, 4], "col3": [
    >>> df.shape
    (2, 3)
    """
    print('hello from shape method')
    return len(self.index), len(self.columns)
```

# Physical structure of package

- When write your own python package?
- Opinion: only if you *absolutely have to*
- A very simple pkg example
  - <https://github.com/jakekrol/aibou>


📁 aibou	hotfix: add missing __init__.py	2 years ago
📁 screenshots	screenshots of game	2 years ago
📄 .gitignore	pypi-release: aibou 1.0.6	2 years ago
📄 LICENSE	pypi-release: aibou 1.0.6	2 years ago
📄 MANIFEST.in	pypi-release: aibou 1.0.6	2 years ago
📄 README.md	README: grammar	2 years ago
📄 pyproject.toml	pypi-release: aibou 1.0.6	2 years ago

## The Scientific Method as an Ongoing Process



# Physical structure of package

- Pandas is just a directory with python code and package metadata



# Conclusions

- Python libraries can speed up science
- **NumPy**: numeric processing & random number generation
- **Pandas**: “nuanced” tabular operations
- **Matplotlib**: visualization
- Limitations and alternatives for each
- “**Libraries**” (dir) are composed of “**packages**” (dir) which are composed of “**modules**” (files)

